

Markov Chain Monte Carlo for Inverse Problems

David Ochsner

June 15, 2020

Contents

1	Theory	2
1.1	Papers	2
1.1.1	Stuart et al: Inverse Problems: A Bayesian Perspective [3]	2
1.1.2	Cotter et al: MCMC for functions [1]	2
1.1.3	Schneider et al: Earth System Modeling 2.0 [2]	3
1.2	Small results	3
1.2.1	Gaussian in infinite dimensions	3
1.2.2	Bayes' Formula & Radon-Nikodym Derivative	4
1.2.3	Acceptance Probability for Metropolis-Hastings	4
1.2.4	Potential for Bayes'-MCMC when sampling from analytic distributions	5
1.2.5	Acceptance Probabilities for different MCMC Proposers	6
1.2.6	Different formulations of multivariate Gaussians	6
1.2.7	Autocorrelation of non-centered distributions	7
2	Framework/Package Structure	10
2.1	Distributions	10
2.2	Potentials	11
2.2.1	AnalyticPotential	12
2.2.2	EvolutionPotential	12
2.3	Proposers	12
2.3.1	StandardRWProposer	12
2.3.2	pCNProposer	13
2.4	Accepters	13
2.4.1	AnalyticAcceptor	13
2.4.2	StandardRWAaccepter	13

2.4.3	pCNAccepter	14
2.4.4	CountedAcceptor	14
2.5	Sampler	14
3	Results	14
3.1	Analytic sampling from a bimodal Gaussian	14
3.1.1	Setup	14
3.1.2	Result	15
3.2	Bayesian inverse problem for $\mathcal{G}(u) = \langle g, u \rangle$	15
3.3	Bayesian inverse problem for $\mathcal{G}(u) = g(u + \beta u^3)$	17
3.4	Lorenz96 model	18
3.4.1	Model	18
3.4.2	Model implementation	22
3.4.3	MCMC	22
3.5	Perturbed Riemann problem for Burgers' equation	27
3.5.1	Model	27
3.5.2	MCMC	27

1 Theory

1.1 Papers

1.1.1 Stuart et al: Inverse Problems: A Bayesian Perspective [3]

Theoretical Background

Notation Central equation:

$$y = \mathcal{G}(u) + \eta$$

with:

- $y \in \mathbb{R}^q$: data
- $u \in \mathbb{R}^n$: IC ("input to mathematical model")
- $\mathcal{G}(\cdot) : \mathbb{R}^n \rightarrow \mathbb{R}^q$: observation operator
- η : mean zero RV, observational noise (a.s. $\eta \sim \mathcal{N}(0, \mathcal{C})$)

1.1.2 Cotter et al: MCMC for functions [1]

Implementation, MCMC in infinite dimensions

1.1.3 Schneider et al: Earth System Modeling 2.0 [2]

Example for MCMC on ODE

1.2 Small results

1.2.1 Gaussian in infinite dimensions

This section is quite a mess, maybe you could suggest a not-too-technical introduction to infinite dimensional Gaussian measures?

Wiki: Definition of Gaussian measure uses Lebesgue measure. However, the Lebesgue-Measure is not defined in an infinite-dimensional space (wiki).

Can still define a measure to be Gaussian if we demand all push-forward measures via a linear functional onto \mathbb{R} to be a Gaussian. (What about the star (E^*, L^*) in the wiki-article? Are they dual-spaces?) (What would be an example of that? An example for a linear functional on an inf-dims space given on wikipedia is integration. What do we integrate? How does this lead to a Gaussian?)

How does this fit with the description in [1]? -> Karhunen-Loève

What would be an example of a covariance operator in infinite dimensions? The Laplace-Operator operates on functions, the eigenfunctions would be *sin*, *cos* (I think? This might not actually be so easy, see Dirichlet Eigenvalues). Are the eigenvalues square-summable?

Anyway, when a inf-dim Gaussian is given as a KL-Expansion, an example of a linear functional given as $f(u) = \langle \phi_i, u \rangle$ for ϕ_i an eigenfunction of \mathcal{C} , then I can see the push-forward definition of inf-dim Gaussians satisfied. (\mathcal{C} spd, so ϕ_i s are orthogonal, so we just end up with one of the KH-"components" which is given to be $\mathcal{N}(0, 1)$).

The problem is not actually in $\exp(-1/2x^T\mathcal{C}^{-1}x)$. What about $\exp(-1/2\|\mathcal{C}^{-1/2}x\|)$?

What about the terminology in [1]? Absolutely continuous w.r.t a measure for example?

How is the square root of an operator defined? For matrices, there seems to be a freedom in choosing whether $A = BB$ or $A = BB^T$ for $B = A^{1/2}$. The latter definition seems to be more useful when working with Cholesky factorizations (cf. <https://math.stackexchange.com/questions/2767873/why-is-the-square-root-of-cholesky-decomposition-equal-to-the-lower> but for example in the wiki-article about the matrix (operator) square root (https://en.wikipedia.org/wiki/Square_root_of_a_matrix): "The Cholesky factorization provides another particular example of square root, which should not be confused with the unique non-negative square root."

1.2.2 Bayes' Formula & Radon-Nikodym Derivative

Bayes' Formula is stated using the Radon-Nikodym Derivative in both [1] and [3]:

$$\frac{d\mu}{d\mu_0} \propto L(u),$$

where $L(u)$ is the likelihood.

Write the measures as $d\mu = \rho(u)du$ and $d\mu_0 = \rho_0(u)du$ with respect to the standard Lesbesgue measure. Then we have

$$\int f(u)\rho(u)du = \int f(u)d\mu(u) = \int f(u)\frac{d\mu(u)}{d\mu_0(u)}d\mu_0 = \int f(u)\frac{d\mu(u)}{d\mu_0(u)}\rho_0(u)du,$$

provided that $d\mu$, $d\mu_0$ and f are nice enough (which they are since we're working with Gaussians). This holds for all test functions f , so it must hold pointwise:

$$\frac{d\mu(u)}{d\mu_0(u)} = \frac{\rho(u)}{\rho_0(u)}.$$

Using this we recover the more familiar formulation of Bayes' formula:

$$\frac{\rho(u)}{\rho_0(u)} \propto L(u).$$

1.2.3 Acceptance Probability for Metropolis-Hastings

A Markov process with transition probabilities $t(y|x)$ has a stationary distribution $\pi(x)$.

- The existence of $\pi(x)$ follows from *detailed balance*:

$$\pi(x)t(y|x) = \pi(y)t(x|y).$$

Detailed balance is sufficient but not necessary for the existence of a stationary distribution.

- Uniqueness of $\pi(x)$ follows from the Ergodicity of the Markov process. For a Markov process to be Ergodic it has to:
 - not return to the same state in a fixed interval
 - reach every state from every other state in finite time

The Metropolis-Hastings algorithm constructs transition probabilities $t(y|x)$ such that the two conditions above are satisfied and that $\pi(x) = P(x)$, where $P(x)$ is the distribution we want to sample from.

Rewrite detailed balance as

$$\frac{t(y|x)}{t(x|y)} = \frac{P(y)}{P(x)}.$$

Split up the transition probability into proposal $g(y|x)$ and acceptance $a(y, x)$. Then detailed balance requires

$$\frac{a(y, x)}{a(x, y)} = \frac{P(y)g(x|y)}{P(x)g(y|x)}.$$

Choose

$$a(y, x) = \min \left\{ 1, \frac{P(y)g(x|y)}{P(x)g(y|x)} \right\}$$

to ensure that detailed balance is always satisfied. Choose $g(y|x)$ such that ergodicity is fulfilled.

If the proposal is symmetric ($g(y|x) = g(x|y)$), then the acceptance takes the simpler form

$$a(y, x) = \min \left\{ 1, \frac{P(y)}{P(x)} \right\}. \quad (1)$$

Since the target distribution $P(x)$ only appears as a ratio, normalizing factors can be ignored.

1.2.4 Potential for Bayes'-MCMC when sampling from analytic distributions

How can we use formulations of Metropolis-Hastings-MCMC algorithms designed to sample from posteriors when want to sample from probability distribution with an easy analytical expression?

Algorithms for sampling from a posterior sample from

$$\rho(u) \propto \rho_0(u) \exp(-\Phi(u)),$$

where ρ_0 is the prior and $\exp(-\Phi(u))$ is the likelihood. Normally, we have an efficient way to compute the likelihood.

When we have an efficient way to compute the posterior ρ and we want to sample from it, the potential to do that is:

$$\Phi(u) = \ln(\rho_0(u)) - \ln(\rho(u)),$$

where an additive constant from the normalization was omitted since only potential differences are relevant.

When working with a Gaussian prior $\mathcal{N}(0, \mathcal{C})$, the potential takes the form

$$\Phi(u) = -\ln \rho(u) - \frac{1}{2} \left\| \mathcal{C}^{-1/2} u \right\|^2.$$

When inserting this into the acceptance probability for the standard random walk MCMC given in formula (1.2) in [1], the two Gaussian-expressions cancel, as do the logarithm and the exponentiation, leaving the simple acceptance described in 1.

This cancellation does not happen when using the pCN-Acceptance probability. This could explain the poorer performance of pCN when directly sampling a probability distribution.

1.2.5 Acceptance Probabilities for different MCMC Proposers

Start from Bayes' formula and rewrite the likelihood $L(u)$ as $\exp(-\Phi(u))$ for a positive scalar function Φ called the potential:

$$\frac{\rho(u)}{\rho_o(u)} \propto \exp(\Phi(u)).$$

Assuming our prior to be a Gaussian ($\mu_0 \sim \mathcal{N}(0, \mathcal{C})$).

Then

$$\rho(u) \propto \exp \left(-\Phi(u) + \frac{1}{2} \left\| \mathcal{C}^{-1/2} u \right\|^2 \right),$$

since $u^T \mathcal{C}^{-1} u = (\mathcal{C}^{-1/2} u)^T (\mathcal{C}^{-1/2} u) = \langle \mathcal{C}^{-1/2} u, \mathcal{C}^{-1/2} u \rangle = \left\| \mathcal{C}^{-1/2} u \right\|^2$, where in the first equality we used \mathcal{C} being symmetric.

This is formula (1.2) in [1] and is used in the acceptance probability for the standard random walk (see also Acceptance Probability for Metropolis-Hastings)

$\mathcal{C}^{-1/2} u$ makes problems in infinite dimensions.

Todo: Why exactly is the second term (from the prior) cancelled when doing pCN?

1.2.6 Different formulations of multivariate Gaussians

Is an RV $\xi \sim \mathcal{N}(0, C)$ distributed the same as $C^{1/2} \xi_0$, with $\xi_0 \sim \mathcal{N}(0, \mathcal{I})$?

From wikipedia: Affine transformation $Y = c + BX$ for $X \sim \mathcal{N}(\mu, \Sigma)$ is also a Gaussian $Y \sim \mathcal{N}(c + B\mu, B\Sigma B^T)$. In our case $X \sim \mathcal{N}(0, \mathcal{I})$, so

$Y \sim \mathcal{N}(0, C^{1/2} \mathcal{I} C^{1/2 T}) = \mathcal{N}(0, C)$, since the covariance matrix is positive definite, which means its square root is also positive definite and thus symmetric.

On second thought, it also follows straight from the definition:

$\mathbf{X} \sim \mathcal{N}(\mu, \Sigma) \Leftrightarrow \exists \mu \in \mathbb{R}^k, A \in \mathbb{R}^{k \times l} \text{ s.t. } \mathbf{X} = \mu + A\mathbf{Z} \text{ with } \mathbf{Z}_n \sim \mathcal{N}(0, 1) \text{ i.i.d}$

where $\Sigma = AA^T$.

1.2.7 Autocorrelation of non-centered distributions

A common definition of the autocorrelation function of a series $\{X_t\}$ is (cite something here?)

$$R(\tau) = \mathbb{E}[X_t X_{t+\tau}^*], \quad (2)$$

which can be normalized by $\tilde{R}(\tau) = R(\tau)/R(0)$ ¹.

For calculating R of a finite series $\{X_t\}_{t=1}^T$, the series can either be zero-padded or the summation limits adjusted accordingly:

$$R(\tau) = \sum_{t=1}^{T-\tau} X_t X_{t+\tau} \text{ for } \tau < T \quad (3)$$

This seems to be the definition that is used in the function `np.correlate`: `c_{av}[k] = sum_n a[n+k] * conj(v[n])` (where we get autocorrelation for `a=v=x`).

This gives the expected result for uniformly random noise in $[-1, 1]$. However, when shifting the same distribution by a constant factor to get uniformly random noise in $[0, 2]$, the autocorrelation decays approximately linearly: 1.

To see why this happens, split up the signal into its mean plus a mean-zero perturbation: $X_t = \bar{X} + \tilde{X}_t$, where $\bar{X} = \mathbb{E}[X]$ and $\mathbb{E}[\tilde{X}] = 0$. The normalized autocorrelation is then:

¹Since we're only working with real numbers, the complex conjugate in the definition will be dropped from now on.

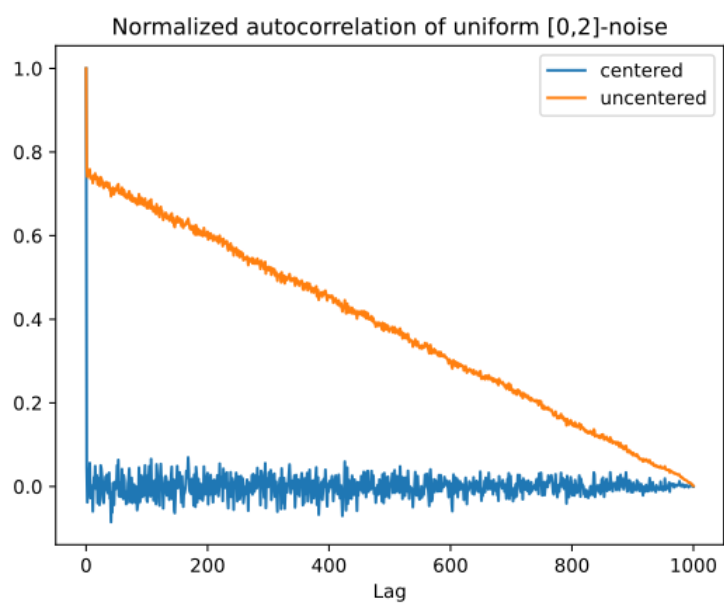


Figure 1: Autocorrelation function of the same signal $\{X_t \sim \mathcal{U}([0, 2])\}$, once computed with `np.correlate` on the original series (uncentered), and once for $\{X_t - 1\}$ (centered).

$$\tilde{R}(\tau) = \frac{\sum_{t=1}^{T-\tau} X_t X_{t+\tau}}{\sum_{t=1}^T X_t^2} \quad (4)$$

$$= \frac{\sum_{t=1}^{T-\tau} (\bar{X} + \tilde{X}_t)(\bar{X} + \tilde{X}_{t+\tau})}{\sum_{t=1}^T (\bar{X} + \tilde{X}_t)^2} \quad (5)$$

$$= \frac{(T-\tau)\bar{X}^2 + \bar{X} \sum_{t=1}^{T-\tau} (\tilde{X}_t + \tilde{X}_{t+\tau}) + \sum_{t=1}^{T-\tau} \tilde{X}_t \tilde{X}_{t+\tau}}{T\bar{X}^2 + 2\bar{X} \sum_{t=1}^T \tilde{X}_t + \sum_{t=1}^T \tilde{X}_t^2} \quad (6)$$

$$= \frac{(T-\tau)\bar{X}^2}{T(\bar{X}^2 + \text{var}(X))}, \quad (7)$$

where the last equality holds because

- $|\sum_{t=1}^{T-\tau} (\tilde{X}_t + \tilde{X}_{t+\tau})| < \epsilon$ and $|\sum_{t=1}^T \tilde{X}_t| < \epsilon$ when $T \gg \tau$ by the weak law of large numbers
- $\sum_{t=1}^{T-\tau} \tilde{X}_t \tilde{X}_{t+\tau} = R(\tau) = 0$ for $\tau \neq 0$ and X "uncorrelated"
- $\sum_{t=1}^T \tilde{X}_t^2 = T \cdot \text{var}(\tilde{X}) = T \cdot \text{var}(X)$

This is a linear function in τ , which is what we see in the plots (plus quite some noise).

For $\bar{X} \gg \text{var}(X)$, we get $\tilde{R}(\tau) = 1 - \tau/T$, which explains nicely why in the "uncentered" autocorrelation always linearly decays to 0, independently of the signal length T .

Considering these points, the python-function that computes the autocorrelation we're actually interested in ² looks like this:

```
def autocorr(x):
    x_centered = x - np.mean(x)
    result = np.correlate(x_centered, x_centered, mode='full')
    # numpy computes the correlation from -\infty to +\infty
    result = result[-len(x):]
    # normalize the result
    return result / result[0]
```

Much of this hassle could be avoided when the expectation value in the definition of the autocorrelation ² would be computed correctly (not just "a

²The function I'm describing here is called auto-covariance function $K_{XX} = \mathbb{E}[(X_t - \mu)(X_{t+\tau} - \mu)^*]$.

posteriori" in the normalizing step with a much too large factor for bigger values of τ). However, this is not possible while still taking advantage of the huge speedup of doing the convolution operation in Fourier space.

2 Framework/Package Structure

The framework is designed to support an easy use case:

```
proposer = StandardRWProposer(beta=0.25, dims=1)
accepter = AnalyticAcceptor(my_distribution)
rng = np.random.default_rng(42)
sampler = MCMCSampler(rw_proposer, accepter, rng)

samples = sampler.run(x_0=0, n_samples=1000)
```

There is only one source of randomness, shared among all classes and supplied by the user. This facilitates reproducibility.

Tests are done with `pytest`.

2.1 Distributions

A class for implementing probability distributions.

```
class DistributionBase(ABC):
    @abstractmethod
    def sample(self, rng):
        """Return a point sampled from this distribution"""
        ...
```

The most important realisation is the `GaussianDistribution`, used in the proposers.

```
class GaussianDistribution(DistributionBase):
    def __init__(self, mean=0, covariance=1):
        ...

    def sample(self, rng):
        ...

    def apply_covariance(self, x):
        ...
```

```

def apply_sqrt_covariance(self, x):
    ...

def apply_precision(self, x):
    ...

def apply_sqrt_precision(self, x):
    ...

```

The design of this class is based on the implementation in `muq2`. The `precision` / `sqrt_precision` is implemented through a Cholesky decomposition, computed in the constructor. This makes applying them pretty fast ($\mathcal{O}(n^2)$).

At the moment there is one class for both scalar and multivariate Gaussians. This introduces some overhead as it has to work with both `float` and `np.array`. Maybe two separate classes would be better.

Also, maybe there is a need to implement a Gaussian using the Karhunen-Loève-Expansion?

2.2 Potentials

A class for implementing the potential resulting from rewriting the likelihood as

$$L(u) = \exp(-\Phi(u)).$$

```

class PotentialBase(ABC):
    """
        Potential used to express the likelihood;
        d mu(u; y) / d mu_0(u) \propto L(u; y)
        Write L(u; y) as exp(-potential(u; y))
    """
    @abstractmethod
    def __call__(self, u):
        ...

    @abstractmethod
    def exp_minus_potential(self, u):
        ...

```

The two functions return $\Phi(u)$ and $\exp(-\Phi(u))$ respectively. Depending on the concrete potential, one or the other is easier to compute.

Potentials are used in the accepters to decide the relative weight of different configurations. They use the `__call__`-method to do that. Especially for high-dimensional error-terms, the value of the pdf of the error term can become very small, so it is important to implement this computing the log-pdf directly instead of manually exponentiating and running into issues with floating point number limitations.

2.2.1 AnalyticPotential

This potential is used when sampling from an analytically computable probability distribution, i.e. a known posterior. In this case

$$\exp(-\Phi(u)) = \frac{\rho(u)}{\rho_0(u)},$$

see theory.org

2.2.2 EvolutionPotential

This potential results when sampling from the model-equation

$$y = \mathcal{G}(u) + \eta,$$

with $\eta \sim \rho$. The resulting potential can be computed as

$$\exp(-\Phi(u)) = \rho(y - \mathcal{G}(u)).$$

2.3 Proposers

Propose a new state v based on the current one u .

```
class ProposerBase(ABC):
    @abstractmethod
    def __call__(self, u, rng):
        ...
```

2.3.1 StandardRWProposer

Propose a new state as

$$v = u + \sqrt{2\delta}\xi,$$

with either $\xi \sim \mathcal{N}(0, \mathcal{I})$ or $\xi \sim \mathcal{N}(0, \mathcal{C})$ (see section 4.2 in [1]).

This leads to a well-defined algorithm in finite dimensions. This is not the case when working on functions (as described in section 6.3 in [1])

2.3.2 pCNProposer

Propose a new state as

$$v = \sqrt{1 - \beta^2}u + \beta\xi,$$

with $\xi \sim \mathcal{N}(0, \mathcal{C})$ and $\beta = \frac{8\delta}{(2+\delta)^2} \in [0, 1]$ (see formula (4.8) in [1]).

This approach leads to an improved algorithm (quicker decorrelation in finite dimensions, nicer properties for infinite dimensions)(see sections 6.2 + 6.3 in [1]).

The wikipedia-article on the Cholesky-factorization mentions the use-case of obtaining a correlated sample from an uncorrelated one by the Cholesky-factor. This is not implemented here.

2.4 Accepters

Given a current state u and a proposed state v , decide if the new state is accepted or rejected.

For sampling from a distribution $P(x)$, the acceptance probability for a symmetric proposal is $a = \min\{1, \frac{P(v)}{P(u)}\}$ (see theory.org)

```
class ProbabilisticAcceptor(AcceptorBase):
    def __call__(self, u, v, rng):
        """Return True if v is accepted"""
        a = self.accept_probability(u, v)
        return a > rng.random()

    @abstractmethod
    def accept_probability(self, u, v):
        ...
```

2.4.1 AnalyticAcceptor

Used when there is an analytic expression of the desired distribution.

```
class AnalyticAcceptor(ProbabilisticAcceptor):
    def accept_probability(self, u, v):
        return self.rho(v) / self.rho(u)
```

2.4.2 StandardRWAcceptor

Based on formula (1.2) in [1]:

$$a = \min\{1, \exp(I(u) - I(v))\},$$

with

$$I(u) = \Phi(u) + \frac{1}{2} \left\| C^{-1/2} u \right\|^2$$

.

See also `theory.org`.

2.4.3 pCNAccepter

Works together with the pCNProposer to achieve the simpler expression for the acceptance

$$a = \min\{1, \exp(\Phi(u) - \Phi(v))\}.$$

2.4.4 CountedAccepter

Stores and forwards calls to an "actual" acceptor. Counts calls and accepts and is used for calculating the acceptance ratio.

2.5 Sampler

The structure of the sampler is quite simple, since it can rely heavily on the functionality provided by the Proposers and Accepters.

```
class MCMCSampler:
    def __init__(self, proposal, acceptance, rng):
        ...

    def run(self, u_0, n_samples, burn_in=1000, sample_interval=200):
        ...

    def _step(self, u, rng):
        ...
```

3 Results

3.1 Analytic sampling from a bimodal Gaussian

3.1.1 Setup

Attempting to recreate the "Computational Illustration" from [1]. They use, among other algorithms, pCN to sample from a 1-D bimodal Gaussian

$$\rho \propto (\mathcal{N}(3, 1) + \mathcal{N}(-3, 1)) \mathbb{1}_{[-10, 10]}.$$

Since the density estimation framework for a known distribution is not quite clear to me from the paper, I don't expect to perfectly replicate their results.

They use a formulation of the prior based on the Karhunen-Lo  ve Expansion that doesn't make sense to me in the 1-D setting (how do I sum infinite eigenfunctions of a scalar?).

The potential for density estimation described in section is also not clear to me (maybe for a similar reason? What is u in the density estimate case?).

I ended up using a normal $\mathcal{N}(0, 1)$ as a prior and the potential described before, and compared the following samplers:

- (1) `StandardRWProposer` ($\delta = 0.25$) + `AnalyticAcceptor`
- (2) `StandardRWProposer` ($\delta = 0.25$) + `StandardRWAaccepter`
- (3) `pCNProposer` ($\beta = 0.25$) + `pCNAcceptor`

The code is in `analytic.py`.

3.1.2 Result

All three samplers are able to reproduce the target density 2




Figure 2: Burn-in: 1000, sample-interval: 200, samples: 500

The autocorrelation decays for all samplers: 3. However, the pCN doesn't do nearly as well as expected. This could be the consequence of the awkward formulation of the potential or a bad prior.

3.2 Bayesian inverse problem for $\mathcal{G}(u) = \langle g, u \rangle$

For $\mathcal{G}(u) = \langle g, u \rangle$ the resulting posterior under a Gaussian prior is again a Gaussian. The model equation is

$$y = \mathcal{G}(u) + \eta$$

with:




Figure 3: AC of bimodal distribution. pCN takes forever to decorrelate

- $y \in \mathbb{R}$
- $u \in \mathbb{R}^n$
- $\eta \sim \mathcal{N}(0, \gamma^2)$ for $\gamma \in \mathbb{R}$

A concrete realization with scalar u :

- $u = 2$
- $g = 3$
- $\gamma = 0.5$
- $y = 6.172$
- prior $\mathcal{N}(0, \Sigma_0 = 1)$

leads to a posterior with mean $\mu = \frac{(\Sigma_0 g) y}{\gamma^2 + \langle g, \Sigma_0 g \rangle} \approx 2$, which is what we see when we plot the result 4. The pCN-Sampler with $\beta = 0.25$ had an acceptance rate of 0.567.




Figure 4: $N = 5000, \mu \approx 2$

For $n > 2$, the resulting posterior can not be plotted anymore. However, it is still Gaussian with given mean & covariance. Can just compare the analytical values to the sample values. Verify that the error decays like $\frac{1}{\sqrt{N}}$.

3.3 Bayesian inverse problem for $\mathcal{G}(u) = g(u + \beta u^3)$

Since the observation operator is not linear anymore, the resulting posterior is not Gaussian in general. However, since the dimension of the input u is 1, it can still be plotted.

The concrete realization with:

- $g = [3, 1]$
- $u = 0.5$
- $\beta = 0$
- $y = [1.672, 0.91]$
- $\gamma = 0.5$

- $\eta \sim \mathcal{N}(0, \gamma^2 I)$
- prior $\mathcal{N}(0, \Sigma_0 = 1)$

however leads to a Gaussian thanks to $\beta = 0$. The mean is $\mu = \frac{\langle g, y \rangle}{\gamma^2 + |g|^2} \approx 0.58$. Plot: 5

The pCN-Sampler with $\beta = 0.25$ (different beta) had an acceptance rate of 0.576.




Figure 5: $N = 5000, \mu \approx 0.58$

For $\beta \neq 0$, the resulting posterior is not a Gaussian. Still $n = 1$, so it can be plotted. Just numerically normalize the analytical expression of the posterior?

3.4 Lorenz96 model

3.4.1 Model

Based on: Properly cite this!

Lorenz, E. N. (1996). Predictability—A problem partly solved. In Reprinted in T. N. Palmer & R. Hagedorn (Eds.), Proceedings Seminar on

Predictability, Predictability of Weather and Climate, Cambridge UP (2006) (Vol. 1, pp. 1–18). Reading, Berkshire, UK: ECMWF.

Equation A system of ODEs, representing the coupling between slow variables X and fast, subgrid variables Y . The system is used in [2] to illustrate different algorithms for earth system modelling.

$$\frac{dX_k}{dt} = -X_{k-1}(X_{k-2} - X_{k+1}) - X_k + F - hc\bar{Y}_k \quad (8)$$

$$\frac{1}{c} \frac{dY_{j,k}}{dt} = -bY_{j+1,k}(Y_{j+2,k} - Y_{j-1,k}) - Y_{j,k} + \frac{h}{J}X_k \quad (9)$$

- $X = [X_0, \dots, X_{K-1}] \in \mathbb{R}^K$
- $Y = [Y_{j,0} | \dots | Y_{j,K-1}] \in \mathbb{R}^{J \times K}$
 $Y_{j,k} = [Y_{0,k}, \dots, Y_{J-1,k}] \in \mathbb{R}^J$
- $\bar{Y}_k = \frac{1}{J} \sum_j Y_{j,k}$
- periodic: $X_K = X_0, Y_{J,k} = Y_{0,k}$
- Parameters $\Theta = [F, h, c, b]$
- h : coupling strength
- c : relative damping
- F : external forcing of the slow variables (large scale forcing)
- b : scale of non-linear interaction of fast variables
- $t = 1 \Leftrightarrow 1$ day (simulation duration is given in days)

b or J ? In the original paper, the equations are given in a different form, namely all explicit occurrences of J above (in the fast-slow interaction) are replaced by b . Since in both concrete realizations (1996 & 2017) are identical and conveniently have $b = J = 10$, the difference doesn't lead to different results for that setup.

"Looking ahead" vs. "Looking back" Comparing nonlinearity terms

$$\begin{aligned} & -X_{k-1}(X_{k-2} - X_{k+1}) \\ & -bY_{j+1,k}(Y_{j+2,k} - Y_{j-1,k}) \end{aligned}$$

for a given Y_k , does the "direction" of the $Z_{k\pm 1}Z_{k\pm 2}$ ($Z = X, Y$) matter?

I don't think so, since the interaction with the other variable is only via point-value and average, and the nonlinearity is periodic.

A bit more formally: The PDE is invariant under "reversing" of the numbering: $Y_{j,k} \rightarrow Y_{J-j,k}$ which is the same as switching $+ \leftrightarrow -$ in the only "asymmetric" term.

Addendum 2 days later: Need to define more clearly what it means for the direction to matter. In the original paper on page 12, it is described how "active areas [...] propagate slowly eastward", while "convective activity tends to propagate westward within the active areas" (rephrased from paper). The paper also explicitly mentions the signs of the subscripts in that context. So some characteristics of the solution are definitely affected. What about the stuff we care about (statistical properties, chaotic behaviour)?

Properties For $K = 36$, $J = 10$ and $\Theta = [F, h, c, b] = [10, 1, 10, 10]$ there is chaotic behaviour.

Energy The nonlinearities conserve the energies within a subsystem:

- $E_X = \sum_k X_k^2$
 $\frac{1}{2} \frac{d(\sum_k X_k^2)}{dt} = \sum_k X_k \frac{dX_k}{dt} = -\sum_k (X_k X_{k-1} X_{k-2} - X_{k-1} X_k X_{k+1}) = 0$,
 where the last equality follows from telescoping + periodicity
- $E_{Y_k} = \sum_j Y_{j,k}^2$
 which follows analogously to the X -case

The interaction between fast and slow variables conserves the total energy:

- $E_T = \sum_k (X_k^2 + \sum_j Y_{j,k}^2)$
 $\frac{1}{2} \frac{dE_T}{dt} = \sum_k X_k \frac{dX_k}{dt} + \sum_j Y_{j,k} \frac{dY_{j,k}}{dt} = \sum_k X_k (-\frac{hc}{J} \sum_j Y_{j,k}) + \sum_j Y_{j,k} (\frac{hc}{J} X_k) =$
 $\sum_k -\frac{hc}{J} X_k (\sum_j Y_{j,k} + \frac{hc}{J} X_k (\sum_j Y_{j,k})) = 0$

In the statistical steady state, the external forcing F (as long as its positive) balances the damping of the linear terms.

Averaged quantities

$$\langle \phi \rangle = \frac{1}{T} \int_{t_0}^{t_0+T} \phi(t) dt$$

(or a sum over discrete values)

Long-term time-average in the statistical steady state: $\langle \cdot \rangle_\infty$

•

$$\langle X_k^2 \rangle_\infty = F \langle X_k \rangle_\infty - hc \langle X_k \bar{Y}_k \rangle_\infty \forall k \quad (10)$$

(multiply X -equation by X , all X_k s are statistically equivalent, $\frac{d\langle X \rangle}{dt} = 0$ in steady state)

•

$$\langle \bar{Y}_k^2 \rangle_\infty = \frac{h}{J} \langle X_k \bar{Y}_k \rangle_\infty \forall k \quad (11)$$

(Quasi) Ergodicity Whether chaotic regions of the phase space of a system are ergodic seems not be an easy question to answer (citation needed probably) ³ Are there any inaccessible regions in phase space for the Lorenz system? I can't think of any. However, there seem to be "traps" that take the system out of it's chaotic behaviour ($X_i = c$, $Y_i = a$). These somehow destroy ergodicity. Are they somehow "measure 0" or something?/. However, for the purposes of this section (which deals with finite time anyway), it is enough to assert that

for the Lorenz system, for sufficiently long times, the time-average converges to the "space-average" over phase-space:

$$\langle f \rangle_\infty = \lim_{T \rightarrow \infty} \int_0^T f(Z(t)) dt = \int_{\mathbb{R}^{K(J+1)}} f(x) \rho(x) dx \quad (12)$$

where $Z(t)$ is a phase space trajectory of the system and $\rho(x)$ is the probability of the system in the statistical steady state to be in state x .

One sufficiently long simulation of the system gives information about all accessible ⁴ initial conditions. As a consequence, as long as the integration time of the system is "long enough", the chosen initial condition is meaningless and can even vary without changing the behaviour of the observation operator.

³Are there any inaccessible regions in phase space for the Lorenz system? I can't think of any. However, there seem to be "traps" that take the system out of it's chaotic behaviour ($X_i = c$, $Y_i = a$). These destroy ergodicity. Are they somehow "measure 0" or something?

⁴Here a more precise definition of ergodicity of the system would help out. What I mean is "all sensible initial conditions".

3.4.2 Model implementation

Implementing the model in python and using a locally 5-th order RK solver yields the following results (initial conditions are just uniformly random numbers in $[0, 1)$ since they don't matter for the long-term evolution of the chaotic system):

Reproducing the results of the original paper Running the setup with $K = 36$, $J = 10$, $(F, h, c, b) = (10, 1, 10, 10)$ gives the following states 6, which qualitatively agree with the results from Lorenz.




Figure 6: System around $T = 20$

The decay of the linear term and the forcing of the slow variables balance out after reaching the steady state, however there is a much bigger fluctuation in the energy than expected 7.

Equilibrium averages Analysis suggests certain long-term averages to be equal in the equilibrium.

3.4.3 MCMC

General point: The RK45 method uses a predictor/corrector step and thus does non-uniform timesteps. However, in the following I compute time-averages with a simple `np.mean`, ignoring the different length of timesteps. It would be not impossible to write my own `time_average(y, t)`-function that takes the non-uniform timesteps into account. However I'm not sure how




Figure 7: Energies in the system. $E_X \gg E_{Y_k} > 0$




Figure 8: RMSE for long-term averages 10 and 11. Averaged over 10 runs

necessary this is, considering a forward-integration takes (2000) timesteps, so I suspect that differences are washed out a bit?

Setup Denote the Lorenz-96 system 8, 9 with parameters $\tilde{u} = [F, h, c, b]$ as $\mathcal{M}[\tilde{u}]$. It acts on the initial condition $z_0 = [X_0, Y_0] \in \mathbb{R}^{K(J+1)}$ to evolve the system for N_t timesteps and generate the phase space trajectory $Z = [X_1 | \dots | X_{N_t}] \in \mathbb{R}^{K(J+1) \times N_t}$.

$$Z = \mathcal{M}[\tilde{u}]z_0$$

Define the "moment function" $f(z) : \mathbb{R}^{K(J+1)} \rightarrow \mathbb{R}^{5K}$

$$f(z) = \begin{bmatrix} X \\ \bar{Y} \\ X^2 \\ X\bar{Y} \\ \bar{Y}^2 \end{bmatrix} \quad (13)$$

The MCMC-Algorithm then samples based on:

$$\langle f \rangle_\infty = \langle f \rangle_T(u) + \eta$$

with:

- $\langle f \rangle_\infty \approx \langle f \rangle_{T_r}$ with $T_r \gg T$ over a simulation $\mathcal{M}[u^*]z_0$
- $\langle f \rangle_T(u)$ the time average over a simulation $\mathcal{M}[u_p + u]z_0$
- Due to the ergodic properties of \mathcal{M} 3.4.1 , it doesn't really matter what z_0 is
- The parameter vector comes in many different variations:
 - $u^* \in \mathbb{R}_{\geq 0}^4$: true underlying parameters, used to compute the "data"
 - $u_p \in \mathbb{R}_{\geq 0}^4$: mean of the prior
 - $u \in \mathbb{R}^4$: perturbations to the prior mean, the actual input to the observation operator
- $\eta \sim \mathcal{N}(0, \Sigma)$, where $\Sigma = r^2[\text{var}(f)_{T_r}]$, where $r \in \mathbb{R}$ is the "noise level"

1. The parameter vector u

The theoretical background assumes the prior to be a centered Gaussian ($\mu = 0$). Specifically, it matters during the proposal step, where the step is taken either with scaled sample from the prior or from a centered Gaussian with the covariance of the prior. A compromise would be to just ignore a nonzero prior-mean in the proposal, however I'm not sure if such a prior has other effects that invalidate the algorithm.

2. "Noise level"

r is scaling of covariance matrix of noise term. This in turn is just step-width in proposal.

TODO: Verify by checking acceptance rate for different noise levels

Concrete parameters The MCMC-Simulation was carried out with the following parameters:

- $K = 6, J = 4$
- Reference Simulation to get $\langle f \rangle_\infty$ and Σ :
 - $u^* = [F^*, h^*, c^*, b^*] = [10, 10, 1, 10]$
 - $T_r = 500$
- Noise $\eta \sim \mathcal{N}(0, \Sigma)$ with $\Sigma = r^2 \text{diag}(\text{var}(f_{T_r})) \in \mathbb{R}^{5K \times 5K}$
- Noise level $r = 0.5$
- Prior $\mathcal{N}(u_p, \Sigma_0)$
 - $u_p = [F_0, h_0, b_0] = [12, 8, 9]$
 - $\Sigma_0 = \text{diag}([10, 1, 10])$
 - c was excluded from the sampling since it is very hard to estimate ("bad statistics")
 - The prior was chosen closer to the true value to make the job of the algorithm easier
- Sampling with pCN proposer and acceptor with $\beta = 0.25$
 - Evaluating the observation operator with a model-simulation of $T = 20$

- Start sampling very close to true value: $u_0 = [-1.9, 1.9, 0.9]$ so that $u^* \approx u_p + u_0$
- This means we can use a short burn-in of 100
- Sample $N = 2500$ with a sample-interval of 2
- The sample interval of 2 is very short, especially considering the long correlation time see below. But 2 is also what they used in the ESM paper.

Result

Density plot for posterior The resulting density plots show a improvement from the prior towards the true value 9. The estimation of the parameter F seems to be easier than b , where the prior and the posterior seem pretty much identical.

This slight improvement is however not unexpected, as the simulations I've done are much shorter than the ones in [2] $(K, J) = (6, 4)$ vs $(36, 10)$, $T = 20$ vs 100 , $T_r = 500$ vs $46,416$)

Should I do some more analysis here, like reporting sample means and covariances to compare posterior/prior not just visually?

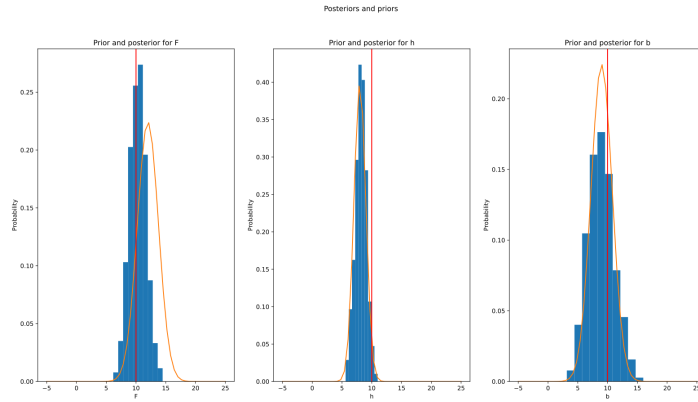


Figure 9: Prior and Posteriors after a 5000 sample MCMC run

ACF The autocorrelation decays for all three variables. As expected from the accuracy of the posteriors, the autocorrelation of F decays much faster than that of b . This simulation was done with a value of $\beta = 0.5$,

which controls the "step size" of the proposer, and resulted in an acceptance rate of around 0.6. The value for β can now be tuned in such a way to get the fastest decay of the autocorrelation, which happens when the steps taken during sampling are big enough to quickly decorrelate the chain, while not being so big that the acceptor declines too many of the steps.

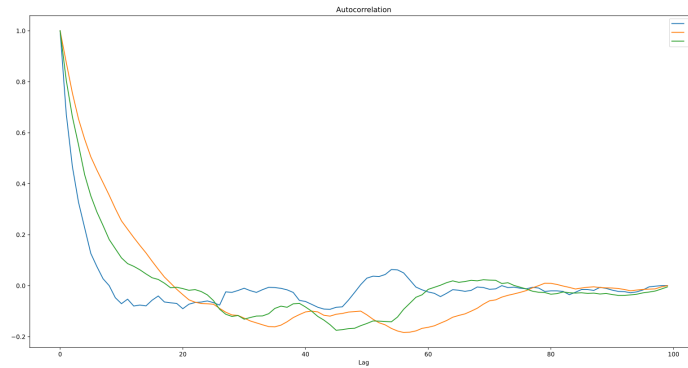


Figure 10: Autocorrelation of during the MCMC sampling. The functions are averaged over ten distinct parts of the chain

3.5 Perturbed Riemann problem for Burgers' equation

3.5.1 Model

Burgers' Equation Blabla

Riemann Problem Blabla

Rusanov FVM Blabla

3.5.2 MCMC

Setup

Bayesian Inversion for Burgers' equation: Test case

We consider the Burgers' equation,

$$(1) \quad u_t + \left(\frac{u^2}{2} \right)_x = 0,$$

in the domain $x \in [-1, 1]$ and $t \in [0, 1]$. The continuous problem will be discretized with the Rusanov scheme.

We consider (1) with the following initial conditions,

$$(2) \quad \begin{aligned} u(x, 0) &= 1 + \delta_1, & \text{if } x < \sigma, \\ &= \delta_2, & \text{if } x > \sigma, \end{aligned}$$

Here, $\delta_1, \delta_2, \sigma \in [-0.05, 0.05]$.

As observables, we will choose the following,

$$(3) \quad L_i = 10 \int_{x_i - 0.05}^{x_i + 0.05} u(x, 1) dx.$$

for $1 \leq i \leq 4$, with $x_1 = -0.5$, $x_2 = -0.25$, $x_3 = 0.25$ and $x_4 = 0.5$.

For the observational noise η_i , with $1 \leq i \leq 4$, we choose $\eta_i \in \mathcal{N}(0, \beta)$ with $\beta = 0.05$ and with each noise variable being uncorrelated from the other.

The ground truth i.e, the underlying measure for the three parameters is given by,

$$(4) \quad \nu_{\delta_1} = \delta_{0.025}, \quad \nu_{\delta_2} = \delta_{-0.025}, \quad \nu_{\sigma} = \delta_{-0.02}.$$

Here δ_a denotes the Dirac measure concentrated at a i.e, the ground truth is deterministic.

As priors, we can start with Gaussian priors of the following type,

$$(5) \quad \nu_{\delta_1}^0 = \mathcal{N}(1.5, \gamma), \quad \nu_{\delta_2}^0 = \mathcal{N}(0.25, \gamma), \quad \nu_{\sigma}^0 = \mathcal{N}(1.3, \gamma),$$

with $\gamma = 0.25$.

Result The evolution of the Markov chain is shown in this figure: 12

The autocorrelation in decaying not as fast as I could have hoped for 3.5.2.

The resulting posteriors look like this: 13. It is apparent from the investigation of the Markov chain and the autocorrelation that a burn-in of 100 and

a sample interval of 5 are both too short to get uncorrelated samples from the steady state. Especially the too short burn-in is visible in the posteriors.





Figure 11: Setup for the MCMC experiment. The values for u at $T = 1$, once for the unperturbed Riemann problem, once for the ground truth of the simulation. The green rectangles are the measurement intervals of the observation operator : $\int_{x_i-0.05}^{x_i+0.05} u(x, 1)dx$, $x_i \in \{-0.5, -0.25, 0.25, 0.5, 0.75\}$.



When the priors are much closer to the ground truth, the chain is able to find the minimum faster, resulting in sharper posteriors 14. I guess this is just because this cuts the burn-in time drastically, as we start basically with the correct u .

With the current setup of measurement points and the discontinuous nature of u (see this figure 11), the "objective function value" is also "pretty




Figure 12: Evolution of the Markov chain with $n=1100$. The spinup is also shown here. The red value show the shock location for their respective δ , σ values, and the red intervals are the measurement intervals. The black lines are the underlying ground truth. It is apparent that a quite long spinup is necessary for the chain to arrive at a steady state.

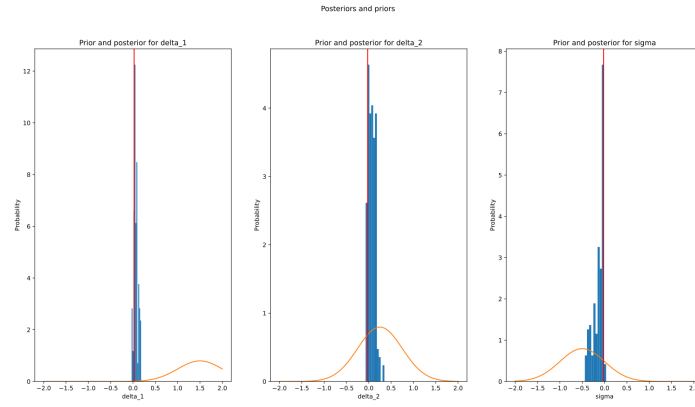


Figure 13: Posteriors of the MCMC-Simulation, obtained from the Chain shown above. $\text{burn}_{\text{in}}=100$, $\text{sample}_{\text{interval}}=5$.




Figure 14: The same setup as before, but the means of the priors are much closer to the ground truth

discontinuous", which gives the Markov chain a hard time moving towards the optimum (interpreting the steps as a stochastic gradient descent).

This manifests itself in the very low acceptance ratio (~ 0.1), easily visible in the later parts of the chain. This results in a pretty bad exploration of the state space around the steady state. This could probably be "fixed" by using priors with smaller variance or by decreasing the β of the MCMC-proposer (are these two equivalent? Not really I think, since a large β also "pulls the proposed states towards 0". This effect would actually be a reason to require priors to be mean-0).

Getting the chain to take smaller steps (at the moment the characteristic stepsize is bigger than the region of the state-space we want to explore ($\beta \cdot \gamma = 0.025$ vs. 0.01 the measurement interval)) would however mean a much much longer burn-in, so the prior-means would have to be chosen closer to the true values.

References

- [1] S. L. Cotter, G. O. Roberts, A. M. Stuart, and D. White. MCMC Methods for Functions: Modifying Old Algorithms to Make Them Faster. *Statistical Science*, 28(3):424–446, August 2013. Publisher: Institute of Mathematical Statistics.

- [2] Tapio Schneider, Shiwei Lan, Andrew Stuart, and João Teixeira. Earth System Modeling 2.0: A Blueprint for Models That Learn From Observations and Targeted High-Resolution Simulations. *Geophysical Research Letters*, 44(24):12,396–12,417, 2017. _eprint: <https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1002/2017GL076101>.
- [3] A. M. Stuart. Inverse problems: A Bayesian perspective. *Acta Numerica*, 19:451–559, May 2010.