

# Markov Chain Monte Carlo for Inverse Problems

David Ochsner

May 6, 2020

## Contents

<b>1</b>	<b>Theory</b>	<b>2</b>
1.1	Papers . . . . .	2
1.1.1	Stuart et al: Inverse Problems: A Bayesian Perspective [3] . . . . .	2
1.1.2	Cotter et al: MCMC for functions [1] . . . . .	2
1.1.3	Schneider et al: Earth System Modeling 2.0 [2] . . . . .	2
1.2	Small results . . . . .	2
1.2.1	Gaussian in infinite dimensions . . . . .	2
1.2.2	Bayes' Formula & Radon-Nikodym Derivative . . . . .	3
1.2.3	Acceptance Probability for Metropolis-Hastings . . . . .	4
1.2.4	Potential for Bayes'-MCMC when sampling from analytic distributions . . . . .	5
1.2.5	Acceptance Probabilities for different MCMC Proposers . . . . .	6
1.2.6	Different formulations of multivariate Gaussians . . . . .	6
<b>2</b>	<b>Implementation</b>	<b>6</b>
2.1	Framework/Package Structure . . . . .	6
2.1.1	Distributions . . . . .	7
2.1.2	Potentials . . . . .	8
2.1.3	Proposers . . . . .	9
2.1.4	Accepters . . . . .	10
2.1.5	Sampler . . . . .	11
2.2	Results . . . . .	11
2.2.1	Analytic sampling from a bimodal Gaussian . . . . .	11
2.2.2	Bayesian inverse problem for $\mathcal{G}(u) = \langle g, u \rangle$ . . . . .	12
2.2.3	Bayesian inverse problem for $\mathcal{G}(u) = g(u + \beta u^3)$ . . . . .	14
2.2.4	Geophysics example: Lorenz-96 model . . . . .	15

# 1 Theory

## 1.1 Papers

### 1.1.1 Stuart et al: Inverse Problems: A Bayesian Perspective [3]

Theoretical Background

1. Notation Central equation:

$$y = \mathcal{G}(u) + \eta$$

with:

- $y \in \mathbb{R}^q$ : data
- $u \in \mathbb{R}^n$ : IC ("input to mathematical model")
- $\mathcal{G}(\cdot) : \mathbb{R}^n \rightarrow \mathbb{R}^q$ : observation operator
- $\eta$ : mean zero RV, observational noise (a.s.  $\eta \sim \mathcal{N}(0, \mathcal{C})$ )

### 1.1.2 Cotter et al: MCMC for functions [1]

Implementation, MCMC in infinite dimensions

### 1.1.3 Schneider et al: Earth System Modeling 2.0 [2]

Example for MCMC on ODE

## 1.2 Small results

### 1.2.1 Gaussian in infinite dimensions

This section is quite a mess, maybe you could suggest a not-too-technical introduction to infinite dimensional Gaussian measures?

Wiki: Definition of Gaussian measure uses Lebesgue measure. However, the Lebesgue-Measure is not defined in an infinite-dimensional space (wiki).

Can still define a measure to be Gaussian if we demand all push-forward measures via a linear functional onto  $\mathbb{R}$  to be a Gaussian. (What about the star ( $E^*$ ,  $L^*$ ) in the wiki-article? Are they dual-spaces?) (What would be an example of that? An example for a linear functional on an inf-dims space given on wikipedia is integration. What do we integrate? How does this lead to a Gaussian?)

How does this fit with the description in [1]? -> Karhunen-Loève

What would be an example of a covariance operator in infinite dimensions? The Laplace-Operator operates on functions, the eigenfunctions would be  $\sin$ ,  $\cos$  (I think? This might not actually be so easy, see Dirichlet Eigenvalues). Are the eigenvalues square-summable?

Anyway, when a inf-dim Gaussian is given as a KL-Expansion, an example of a linear functional given as  $f(u) = \langle \phi_i, u \rangle$  for  $\phi_i$  an eigenfunction of  $\mathcal{C}$ , then I can see the push-forward definition of inf-dim Gaussians satisfied. ( $\mathcal{C}$  spd, so  $\phi_i$  s are orthogonal, so we just end up with one of the KH-"components" which is given to be  $\mathcal{N}(0, 1)$ ).

The problem is not actually in  $\exp(-1/2x^T\mathcal{C}^{-1}x)$ . What about  $\exp(-1/2\|\mathcal{C}^{-1/2}x\|)$ ?

What about the terminology in [1]? Absolutely continuous w.r.t a measure for example?

How is the square root of an operator defined? For matrices, there seems to be a freedom in choosing whether  $A = BB$  or  $A = BB^T$  for  $B = A^{1/2}$ . The latter definition seems to be more useful when working with Cholesky factorizations (cf. <https://math.stackexchange.com/questions/2767873/why-is-the-square-root-of-cholesky-decomposition-equal-to-the-lower> but for example in the wiki-article about the matrix (operator) square root ([https://en.wikipedia.org/wiki/Square\\_root\\_of\\_a\\_matrix](https://en.wikipedia.org/wiki/Square_root_of_a_matrix)): "The Cholesky factorization provides another particular example of square root, which should not be confused with the unique non-negative square root."

### 1.2.2 Bayes' Formula & Radon-Nikodym Derivative

Bayes' Formula is stated using the Radon-Nikodym Derivative in both [1] and [3]:

$$\frac{d\mu}{d\mu_0} \propto L(u),$$

where  $L(u)$  is the likelihood.

Write the measures as  $d\mu = \rho(u)du$  and  $d\mu_0 = \rho_0(u)du$  with respect to the standard Lesbesgue measure. Then we have

$$\int f(u)\rho(u)du = \int f(u)d\mu(u) = \int f(u)\frac{d\mu(u)}{d\mu_0(u)}d\mu_0 = \int f(u)\frac{d\mu(u)}{d\mu_0(u)}\rho_0(u)du,$$

provided that  $d\mu$ ,  $d\mu_0$  and  $f$  are nice enough (which they are since we're working with Gaussians). This holds for all test functions  $f$ , so it must hold pointwise:

$$\frac{d\mu(u)}{d\mu_0(u)} = \frac{\rho(u)}{\rho_0(u)}.$$

Using this we recover the more familiar formulation of Bayes' formula:

$$\frac{\rho(u)}{\rho_o(u)} \propto L(u).$$

### 1.2.3 Acceptance Probability for Metropolis-Hastings

A Markov process with transition probabilities  $t(y|x)$  has a stationary distribution  $\pi(x)$ .

- The existence of  $\pi(x)$  follows from *detailed balance*:

$$\pi(x)t(y|x) = \pi(y)t(x|y).$$

Detailed balance is sufficient but not necessary for the existence of a stationary distribution.

- Uniqueness of  $\pi(x)$  follows from the Ergodicity of the Markov process. For a Markov process to be Ergodic it has to:
  - not return to the same state in a fixed interval
  - reach every state from every other state in finite time

The Metropolis-Hastings algorithm constructs transition probabilities  $t(y|x)$  such that the two conditions above are satisfied and that  $\pi(x) = P(x)$ , where  $P(x)$  is the distribution we want to sample from.

Rewrite detailed balance as

$$\frac{t(y|x)}{t(x|y)} = \frac{P(y)}{P(x)}.$$

Split up the transition probability into proposal  $g(y|x)$  and acceptance  $a(y, x)$ . Then detailed balance requires

$$\frac{a(y, x)}{a(x, y)} = \frac{P(y)g(x|y)}{P(x)g(y|x)}.$$

Choose

$$a(y, x) = \min \left\{ 1, \frac{P(y)g(x|y)}{P(x)g(y|x)} \right\}$$

to ensure that detailed balance is always satisfied. Choose  $g(y|x)$  such that ergodicity is fulfilled.

If the proposal is symmetric ( $g(y|x) = g(x|y)$ ), then the acceptance takes the simpler form

$$a(y, x) = \min \left\{ 1, \frac{P(y)}{P(x)} \right\}. \quad (1)$$

Since the target distribution  $P(x)$  only appears as a ratio, normalizing factors can be ignored.

#### 1.2.4 Potential for Bayes'-MCMC when sampling from analytic distributions

How can we use formulations of Metropolis-Hastings-MCMC algorithms designed to sample from posteriors when want to sample from probability distribution with an easy analytical expression?

Algorithms for sampling from a posterior sample from

$$\rho(u) \propto \rho_0(u) \exp(-\Phi(u)),$$

where  $\rho_0$  is the prior and  $\exp(-\Phi(u))$  is the likelihood. Normally, we have an efficient way to compute the likelihood.

When we have an efficient way to compute the posterior  $\rho$  and we want to sample from it, the potential to do that is:

$$\Phi(u) = \ln(\rho_0(u)) - \ln(\rho(u)),$$

where an additive constant from the normalization was omitted since only potential differences are relevant.

When working with a Gaussian prior  $\mathcal{N}(0, \mathcal{C})$ , the potential takes the form

$$\Phi(u) = -\ln \rho(u) - \frac{1}{2} \left\| \mathcal{C}^{-1/2} u \right\|^2.$$

When inserting this into the acceptance probability for the standard random walk MCMC given in formula (1.2) in [1], the two Gaussian-expressions cancel, as do the logarithm and the exponentiation, leaving the simple acceptance described in 1.

This cancellation does not happen when using the pCN-Acceptance probability. This could explain the poorer performance of pCN when directly sampling a probability distribution.

### 1.2.5 Acceptance Probabilities for different MCMC Proposers

Start from Bayes' formula and rewrite the likelihood  $L(u)$  as  $\exp(-\Phi(u))$  for a positive scalar function  $\Phi$  called the potential:

$$\frac{\rho(u)}{\rho_o(u)} \propto \exp(\Phi(u)).$$

Assuming our prior to be a Gaussian ( $\mu_0 \sim \mathcal{N}(0, C)$ ).

Then

$$\rho(u) \propto \exp\left(-\Phi(u) + \frac{1}{2}\|C^{-1/2}u\|^2\right),$$

since  $u^T C^{-1}u = (C^{-1/2}u)^T(C^{-1/2}u) = \langle C^{-1/2}u, C^{-1/2}u \rangle = \|C^{-1/2}u\|^2$ , where in the first equality we used  $C$  being symmetric.

This is formula (1.2) in [1] and is used in the acceptance probability for the standard random walk (see also Acceptance Probability for Metropolis-Hastings)

$C^{-1/2}u$  makes problems in infinite dimensions.

Todo: Why exactly is the second term (from the prior) cancelled when doing pCN?

### 1.2.6 Different formulations of multivariate Gaussians

Is an RV  $\xi \sim \mathcal{N}(0, C)$  distributed the same as  $C^{1/2}\xi_0$ , with  $\xi_0 \sim \mathcal{N}(0, \mathcal{I})$ ?

From wikipedia: Affine transformation  $Y = c + BX$  for  $X \sim \mathcal{N}(\mu, \Sigma)$  is also a Gaussian  $Y \sim \mathcal{N}(c + B\mu, B\Sigma B^T)$ . In our case  $X \sim \mathcal{N}(0, \mathcal{I})$ , so  $Y \sim \mathcal{N}(0, C^{1/2}\mathcal{I}C^{1/2}) = \mathcal{N}(0, C)$ , since the covariance matrix is positive definite, which means it's square root is also positive definite and thus symmetric.

On second thought, it also follows straight from the definition:

$$\mathbf{X} \sim \mathcal{N}(\mu, \Sigma) \Leftrightarrow \exists \mu \in \mathbb{R}^k, A \in \mathbb{R}^{k \times l} \text{ s.t. } \mathbf{X} = \mu + A\mathbf{Z} \text{ with } \mathbf{Z}_n \sim \mathcal{N}(0, 1) \text{ i.i.d}$$

where  $\Sigma = AA^T$ .

## 2 Implementation

### 2.1 Framework/Package Structure

The framework is designed to support an easy use case:

```

proposer = StandardRWProposer(beta=0.25, dims=1)
accepter = AnalyticAcceptor(my_distribution)
rng = np.random.default_rng(42)
sampler = MCMCSampler(rw_proposer, accepter, rng)

samples = sampler.run(x_0=0, n_samples=1000)

```

There is only one source of randomness, shared among all classes and supplied by the user. This facilitates reproducibility.

Tests are done with `pytest`.

### 2.1.1 Distributions

A class for implementing probability distributions.

```

class DistributionBase(ABC):
    @abstractmethod
    def sample(self, rng):
        """Return a point sampled from this distribution"""
        ...

```

The most important realisation is the `GaussianDistribution`, used in the proposers.

```

class GaussianDistribution(DistributionBase):
    def __init__(self, mean=0, covariance=1):
        ...

    def sample(self, rng):
        ...

    def apply_covariance(self, x):
        ...

    def apply_sqrt_covariance(self, x):
        ...

    def apply_precision(self, x):
        ...

    def apply_sqrt_precision(self, x):
        ...

```

The design of this class is based on the implementation in `muq2`. The `precision / sqrt_precision` is implemented through a Cholesky decomposition, computed in the constructor. This makes applying them pretty fast ( $\mathcal{O}(n^2)$ ).

At the moment there is one class for both scalar and multivariate Gaussians. This introduces some overhead as it has to work with both `float` and `np.array`. Maybe two separate classes would be better.

Also, maybe there is a need to implement a Gaussian using the Karhunen-Loève-Expansion?

### 2.1.2 Potentials

A class for implementing the potential resulting from rewriting the likelihood as

$$L(u) = \exp(-\Phi(u)).$$

```
class PotentialBase(ABC):
    """
        Potential used to express the likelihood;
        d mu(u; y) / d mu_0(u) \propto L(u; y)
        Write L(u; y) as exp(-potential(u; y))
    """
    @abstractmethod
    def __call__(self, u):
        ...

    @abstractmethod
    def exp_minus_potential(self, u):
        ...
```

The two functions return  $\Phi(u)$  and  $\exp(-\Phi(u))$  respectively. Depending on the concrete potential, one or the other is easier to compute.

Potentials are used in the accepters to decide the relative weight of different configurations. There, the `PotentialBase.exp_minus_potential` is used.

#### 1. AnalyticPotential

This potential is used when sampling from an analytically computable probability distribution, i.e. a known posterior. In this case

$$\exp(-\Phi(u)) = \frac{\rho(u)}{\rho_0(u)},$$



see 1.2.4

## 2. EvolutionPotential

This potential results when sampling from the model-equation

$$y = \mathcal{G}(u) + \eta,$$

with  $\eta \sim \rho$ . The resulting potential can be computed as

$$\exp(-\Phi(u)) = \rho(y - \mathcal{G}(u)).$$

### 2.1.3 Proposers

Propose a new state  $v$  based on the current one  $u$ .

```
class ProposerBase(ABC):
    @abstractmethod
    def __call__(self, u, rng):
        ...
```

#### 1. StandardRWProposer

Propose a new state as

$$v = u + \sqrt{2\delta}\xi,$$

with either  $\xi \sim \mathcal{N}(0, \mathcal{I})$  or  $\xi \sim \mathcal{N}(0, \mathcal{C})$  (see section 4.2 in [1]).

This leads to a well-defined algorithm in finite dimensions. This is not the case when working on functions (as described in section 6.3 in [1])

#### 2. pCNProposer

Propose a new state as

$$v = \sqrt{1 - \beta^2}u + \beta\xi,$$

with  $\xi \sim \mathcal{N}(0, \mathcal{C})$  and  $\beta = \frac{8\delta}{(2+\delta)^2} \in [0, 1]$  (see formula (4.8) in [1]).

This approach leads to an improved algorithm (quicker decorrelation in finite dimensions, nicer properties for infinite dimensions)(see sections 6.2 + 6.3 in [1]).

The wikipedia-article on the Cholesky-factorization mentions the use-case of obtaining a correlated sample from an uncorrelated one by the Cholesky-factor. This is not implemented here.

### 2.1.4 Accepters

Given a current state  $u$  and a proposed state  $v$ , decide if the new state is accepted or rejected.

For sampling from a distribution  $P(x)$ , the acceptance probability for a symmetric proposal is  $a = \min\{1, \frac{P(v)}{P(u)}\}$  (see 1.2.3)

```
class ProbabilisticAcceptor(AcceptorBase):
    def __call__(self, u, v, rng):
        """Return True if v is accepted"""
        a = self.accept_probability(u, v)
        return a > rng.random()

    @abstractmethod
    def accept_probability(self, u, v):
        ...
```

#### 1. AnalyticAcceptor

Used when there is an analytic expression of the desired distribution.

```
class AnalyticAcceptor(ProbabilisticAcceptor):
    def accept_probability(self, u, v):
        return self.rho(v) / self.rho(u)
```

#### 2. StandardRWAcceptor

Based on formula (1.2) in [1]:

$$a = \min\{1, \exp(I(u) - I(v))\},$$

with

$$I(u) = \Phi(u) + \frac{1}{2} \left\| \mathcal{C}^{-1/2} u \right\|^2$$

.

See also 1.2.5.

#### 3. pCNAcceptor

Works together with the pCNProposer to achieve the simpler expression for the acceptance

$$a = \min\{1, \exp(\Phi(u) - \Phi(v))\}.$$

#### 4. CountedAcceptor

Stores and forwards calls to an "actual" accepter. Counts calls and accepts and is used for calculating the acceptance ratio.

### 2.1.5 Sampler

The structure of the sampler is quite simple, since it can rely heavily on the functionality provided by the Proposers and Accepters.

```
class MCMCSampler:
    def __init__(self, proposal, acceptance, rng):
        ...

    def run(self, u_0, n_samples, burn_in=1000, sample_interval=200):
        ...

    def _step(self, u, rng):
        ...
```

## 2.2 Results

### 2.2.1 Analytic sampling from a bimodal Gaussian

#### 1. Setup

Attempting to recreate the "Computational Illustration" from [1]. They use, among other algorithms, pCN to sample from a 1-D bimodal Gaussian

$$\rho \propto (\mathcal{N}(3, 1) + \mathcal{N}(-3, 1))\mathbb{1}_{[-10, 10]}.$$

Since the density estimation framework for a known distribution is not quite clear to me from the paper, I don't expect to perfectly replicate their results.

They use a formulation of the prior based on the Karhunen-Loève Expansion that doesn't make sense to me in the 1-D setting (how do I sum infinite eigenfunctions of a scalar?).

The potential for density estimation described in section is also not clear to me (maybe for a similar reason? What is  $u$  in the density estimate case?).

I ended up using a normal  $\mathcal{N}(0, 1)$  as a prior and the potential described before, and compared the following samplers:

- (1) StandardRWProposer ( $\delta = 0.25$ ) + AnalyticAcceptor
- (2) StandardRWProposer ( $\delta = 0.25$ ) + StandardRWAaccepter
- (3) pCNProposer ( $\beta = 0.25$ ) + pCNAcceptor

The code is in `analytic.py`.

## 2. Result

All three samplers are able to reproduce the target density 1 2 2.

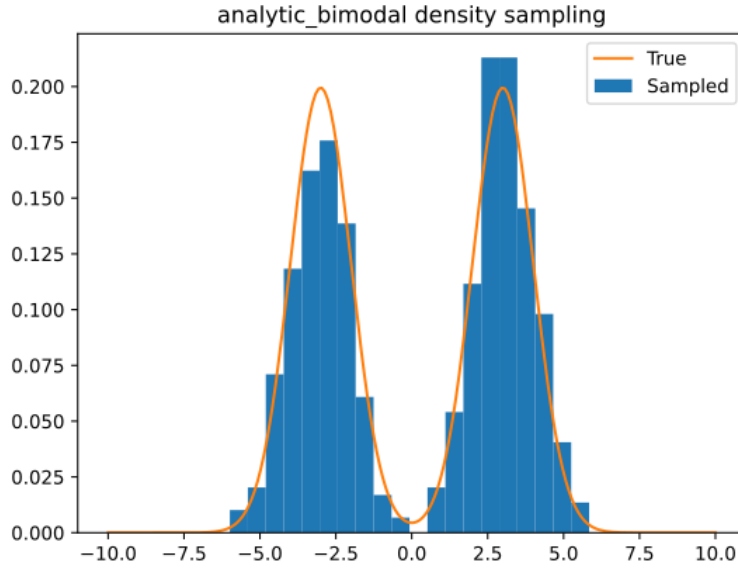


Figure 1: analytic

The autocorrelation decays for all samplers: 4, 5. However, the pCN doesn't do nearly as well as expected. This could be the consequence of the awkward formulation of the potential or a bad prior.

A peculiar thing about the decorrelation of the pCN sampling process is that it somehow is tied to the number of samples, compare 6 and 7. Is this a bug or a misunderstanding of the autocorrelation function?

### 2.2.2 Bayesian inverse problem for $\mathcal{G}(u) = \langle g, u \rangle$

For  $\mathcal{G}(u) = \langle g, u \rangle$  the resulting posterior under a Gaussian prior is again a Gaussian. The model equation is

$$y = \mathcal{G}(u) + \eta$$

with:




Figure 2: standard rw

- $y \in \mathbb{R}$
- $u \in \mathbb{R}^n$
- $\eta \sim \mathcal{N}(0, \gamma^2)$  for  $\gamma \in \mathbb{R}$

A concrete realization with scalar  $u$ :

- $u = 2$
- $g = 3$
- $\gamma = 0.5$
- $y = 6.172$
- prior  $\mathcal{N}(0, \Sigma_0 = 1)$

leads to a posterior with mean  $\mu = \frac{(\Sigma_0 g) y}{\gamma^2 + \langle g, \Sigma_0 g \rangle} \approx 2$ , which is what we see when we plot the result 8. The pCN-Sampler with  $\beta = 0.25$  had an acceptance rate of 0.567.

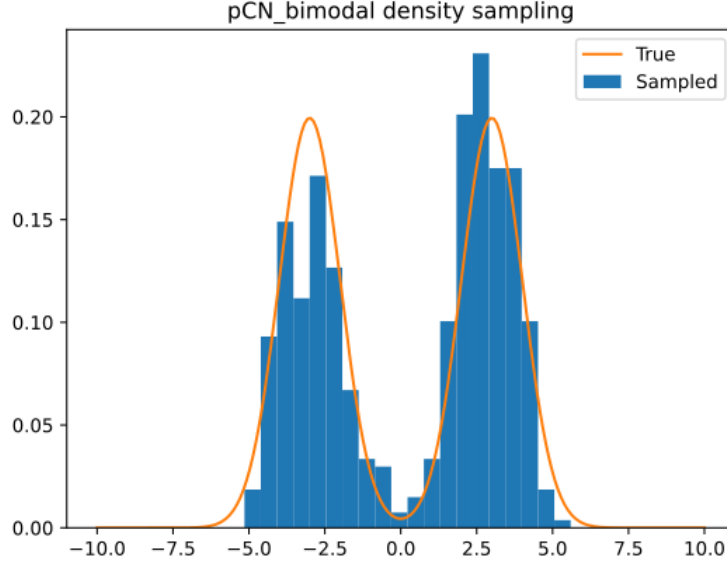


Figure 3: pCN

For  $n > 2$ , the resulting posterior can not be plotted anymore. However, it is still Gaussian with given mean & covariance. Can just compare the analytical values to the sample values. Verify that the error decays like  $\frac{1}{\sqrt{N}}$ .

### 2.2.3 Bayesian inverse problem for $\mathcal{G}(u) = g(u + \beta u^3)$

Since the observation operator is not linear anymore, the resulting posterior is not Gaussian in general. However, since the dimension of the input  $u$  is 1, it can still be plotted.

The concrete realization with:

- $g = [3, 1]$
- $u = 0.5$
- $\beta = 0$
- $y = [1.672, 0.91]$
- $\gamma = 0.5$




Figure 4: AC of standard normal. All samplers decorrelate quickly

- $\eta \sim \mathcal{N}(0, \gamma^2 I)$
- prior  $\mathcal{N}(0, \Sigma_0 = 1)$

however leads to a Gaussian thanks to  $\beta = 0$ . The mean is  $\mu = \frac{\langle g, y \rangle}{\gamma^2 + |g|^2} \approx 0.58$ . Plot: 9

The pCN-Sampler with  $\beta = 0.25$  (different beta) had an acceptance rate of 0.576.

For  $\beta \neq 0$ , the resulting posterior is not a Gaussian. Still  $n = 1$ , so it can be plotted. Just numerically normalize the analytical expression of the posterior?

#### 2.2.4 Geophysics example: Lorenz-96 model

1. Based on:

Lorenz, E. N. (1996). Predictability—A problem partly solved. In Reprinted in T. N. Palmer & R. Hagedorn (Eds.), Proceedings Seminar on Predictability, Predictability of Weather and Climate, Cambridge UP (2006) (Vol. 1, pp. 1–18). Reading, Berkshire, UK: ECMWF.




Figure 5: AC of bimodal distribution. pCN takes forever to decorrelate

## 2. Equation

A system of ODEs, representing the coupling between slow variables  $X$  and fast, subgrid variables  $Y$ .

$$\begin{aligned}\frac{dX_k}{dt} &= -X_{k-1}(X_{k-2} - X_{k+1}) - X_k + F - hc\bar{Y}_k \\ \frac{1}{c} \frac{dY_{j,k}}{dt} &= -bY_{j+1,k}(Y_{j-2,k} - Y_{j+1,k}) - Y_{j,k} + \frac{h}{J}X_k\end{aligned}$$

- $X = [X_0, \dots, X_{K-1}] \in \mathbb{R}^K$
- $Y = [Y_{j,0} | \dots | Y_{j,K-1}] \in \mathbb{R}^{J \times K}$   
 $Y_{j,k} = [Y_{0,k}, \dots, Y_{J-1,k}] \in \mathbb{R}^J$
- $\bar{Y}_k = \frac{1}{J} \sum_j Y_{j,k}$
- periodic:  $X_K = X_0, Y_{J,k} = Y_{0,k}$
- Parameters  $\Theta = [F, h, c, b]$
- $h$ : coupling strength
- $c$ : relative damping






Figure 6: AC of bimodal distribution.

- $F$ : external forcing of the slow variables (large scale forcing)
- $b$ : scale of non-linear interaction of fast variables
- $t = 1 \Leftrightarrow 1$  day (simulation duration is given in days)

### 3. Properties

- For  $K = 36$ ,  $J = 10$  and  $\Theta = [F, h, c, b] = [10, 1, 10, 10]$  there is chaotic behaviour.
- The nonlinearities conserve the energies within a subsystem: (show!)
  - $E_X = \sum_k X_k^2$
  - $E_{Y_k} = \sum_j Y_{j,k}^2$
- The interaction conserves the total energy: (show!)
  - $E_T = \sum_k (X_k^2 + \sum_j Y_{j,k}^2)$
- In the statistical steady state, the external forcing  $F$  (as long as its positive) balances the damping of the linear terms.
- Averaged quantities




Figure 7: AC of bimodal distribution.

- $\langle \phi \rangle = \frac{1}{T} \int_{t_0}^{t_0+T} \phi(t) dt$  (or a sum over discrete values)
- Long-term time-mean in the statistical steady state:  $\langle \cdot \rangle_{\infty}$
- $\langle X^2 \rangle_{\infty} = F \langle X \rangle_{\infty} - hc \langle X \bar{Y} \rangle_{\infty} \quad \forall k$   
(multiply  $X$ -equation by  $X$ , all  $X_k$  s are statistically equivalent,  $\frac{dX}{dt} = 0$  in steady state)
- $\langle \bar{Y}^2 \rangle_{\infty} = \frac{h}{j} \langle X \bar{Y} \rangle_{\infty}$

## References

- [1] S. L. Cotter, G. O. Roberts, A. M. Stuart, and D. White. MCMC Methods for Functions: Modifying Old Algorithms to Make Them Faster. *Statistical Science*, 28(3):424–446, August 2013. Publisher: Institute of Mathematical Statistics.
- [2] Tapio Schneider, Shiwei Lan, Andrew Stuart, and João Teixeira. Earth System Modeling 2.0: A Blueprint for Models That Learn From Observations and Targeted High-Resolution Simulations.




Figure 8:  $N = 5000, \mu \approx 2$

*Geophysical Research Letters*, 44(24):12,396–12,417, 2017. [\\_eprint:  
https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1002/2017GL076101](https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1002/2017GL076101).

- [3] A. M. Stuart. Inverse problems: A Bayesian perspective. *Acta Numerica*, 19:451–559, May 2010.




Figure 9:  $N = 5000, \mu \approx 0.58$