Markov Chain Monte Carlo for Inverse Problems

David Ochsner

May 28, 2020

Contents

1 Theory

1.1 Papers

1.1.1 Stuart et al: Inverse Problems: A Bayesian Perspective [3]

Theoretical Background

Notation   Central equation:

$$y = \mathcal{G}(u) + \eta$$

with:

- $y \in \mathbb{R}^q$: data

- $u \in \mathbb{R}^n$: IC ("input to mathematical model")

- $\mathcal{G}(\cdot) : \mathbb{R}^n \to \mathbb{R}^q$: observation operator

- η: mean zero RV, observational noise (a.s. $\eta \sim \mathcal{N}(0, \mathcal{C})$)

1.1.2 Cotter et al: MCMC for functions [1]

Implementation, MCMC in infinite dimensions

1.1.3 Schneider et al: Earth System Modeling 2.0 [2]

Example for MCMC on ODE

1.2 Small results

1.2.1 Gaussian in infinite dimensions

This section is quite a mess, maybe you could suggest a not-too-technical introduction to infinite dimensional Gaussian measures?

Wiki: Definition of Gaussian measure uses Lesbesgue measure. However, the Lesbesgue-Measure is not defined in an infinite-dimensional space (wiki).

Can still define a measure to be Gaussian if we demand all push-forward measures via a linear functional onto $\mathbb{R}$ to be a Gaussian. (What about the star (E^*, L_*) in the wiki-article? Are they dual-spaces?) (What would be an example of that? An example for a linear functional on an inf-dims space given on wikipedia is integration. What do we integrate? How does this lead to a Gaussian?)

How does this fit with the description in [1]? -> Karhunenen-Loéve

What would be an example of a covariance operator in infinite dimensions? The Laplace-Operator operates on functions, the eigenfunctions would be sin, cos (I think? This might not actually be so easy, see Dirichlet Eigenvalues). Are the eigenvalues square-summable?

Anyway, when a inf-dim Gaussian is given as a KL-Expansion, an example of a linear functional given as $f(u) = \langle \phi_i, u \rangle$ for ϕ_i an eigenfunction of $\mathcal{C}$, then I can see the push-forward definition of inf-dim Gaussians satisfied. ($\mathcal{C}$ spd, so ϕ_i s are orthogonal, so we just end up with one of the KH-"components" which is given to be $\mathcal{N}(0,1)$).

The problem is not actually in $\exp\left(-1/2 x^T \mathcal{C}^{-1} x\right)$. What about $\exp\left(-1/2\|\mathcal{C}^{-1/2}x\|\right)$?

What about the terminology in [1]? Absolutely continuous w.r.t a measure for example?

How is the square root of an operator defined? For matrices, there seems to be a freedom in choosing whether $A = BB$ or $A = BB^T$ for $B = A^{1/2}$. The latter definition seems to be more useful when working with Cholesky factorizations (cf. `https://math.stackexchange.com/questions/2767873/why-is-the-square-root-of-cholesky-decomposition-equal-to-the-lower-` but for example in the wiki-article about the matrix (operator) square root (`https://en.wikipedia.org/wiki/Square_root_of_a_matrix`): "The Cholesky factorization provides another particular example of square root, which should not be confused with the unique non-negative square root."

1.2.2 Bayes' Formula & Radon-Nikodym Derivative

Bayes' Formula is stated using the Radon-Nikodym Derivative in both [1] and [3]:

$$\frac{\mathrm{d}\mu}{\mathrm{d}\mu_0} \propto \mathrm{L}(u),$$

where $\mathrm{L}(u)$ is the likelihood.

Write the measures as $\mathrm{d}\mu = \rho(u)\mathrm{d}u$ and $\mathrm{d}\mu_0 = \rho_0(u)\mathrm{d}u$ with respect to the standard Lesbesgue measure. Then we have

$$\int f(u)\rho(u)\mathrm{d}u = \int f(u)\mathrm{d}\mu(u) = \int f(u)\frac{\mathrm{d}\mu(u)}{\mathrm{d}\mu_0(u)}\mathrm{d}\mu_0 = \int f(u)\frac{\mathrm{d}\mu(u)}{\mathrm{d}\mu_0(u)}\rho_0(u)\mathrm{d}u,$$

provided that $\mathrm{d}\mu$, $\mathrm{d}\mu_0$ and f are nice enough (which they are since we're working with Gaussians). This holds for all test functions f, so it must hold pointwise:

$$\frac{\mathrm{d}\mu(u)}{\mathrm{d}\mu_0(u)} = \frac{\rho(u)}{\rho_o(u)}.$$

Using this we recover the more familiar formulation of Bayes' formula:

$$\frac{\rho(u)}{\rho_o(u)} \propto \mathrm{L}(u).$$

1.2.3 Acceptance Probability for Metropolis-Hastings

A Markov process with transition probabilities $t(y|x)$ has a stationary distribution $\pi(x)$.

- The <u>existence</u> of $\pi(x)$ follows from *detailed balance*:

$$\pi(x)t(y|x) = \pi(y)t(x|y).$$

  Detailed balance is sufficient but not necessary for the existence of a stationary distribution.

- <u>Uniqueness</u> of $\pi(x)$ follows from the Ergodicity of the Markov process. For a Markov processto be Ergodic it has to:

  - not return to the same state in a fixed interval
  - reach every state from every other state in finite time

The Metropolis-Hastings algorithm constructs transition probabilities $t(y|x)$ such that the two conditions above are satisfied and that $\pi(x) = P(x)$, where $P(x)$ is the distribution we want to sample from.

Rewrite detailed balance as

$$\frac{t(y|x)}{t(x|y)} = \frac{P(y)}{P(x)}.$$

Split up the transition probability into proposal $g(y|x)$ and acceptance $a(y,x)$. Then detailed balance requires

$$\frac{a(y,x)}{a(x,y)} = \frac{P(y)g(x|y)}{P(x)g(y|x)}.$$

Choose

$$a(y,x) = \min\left\{1, \frac{P(y)g(x|y)}{P(x)g(y|x)}\right\}$$

to ensure that detailed balance is always satisfied. Choose $g(y|x)$ such that ergodicity is fulfilled.

If the proposal is symmetric $(g(y|x) = g(x|y))$, then the acceptance takes the simpler form

$$a(y,x) = \min\left\{1, \frac{P(y)}{P(x)}\right\}. \tag{1}$$

Since the target distribution $P(x)$ only appears as a ratio, normalizing factors can be ignored.

1.2.4 Potential for Bayes'-MCMC when sampling from analytic distributions

How can we use formulations of Metropolis-Hastings-MCMC algorithms designed to sample from posteriors when want to sample from probability distribution with an easy analytical expression?

Algorithms for sampling from a posterior sample from

$$\rho(u) \propto \rho_0(u)\exp(-\Phi(u)),$$

where ρ_0 is the prior and $\exp(-\Phi(u))$ is the likelihood. Normally, we have an efficient way to compute the likelihood.

When we have an efficient way to compute the posterior ρ and we want to sample from it, the potential to do that is:

$$\Phi(u) = \ln(\rho_0(u)) - \ln(\rho(u)),$$

where an additive constant from the normalization was omitted since only potential differences are relevant.

When working with a Gaussian prior $\mathcal{N}(0, \mathcal{C})$, the potential takes the form

$$\Phi(u) = -\ln\rho(u) - \frac{1}{2}\left\|\mathcal{C}^{-1/2}u\right\|^2.$$

When inserting this into the acceptance probability for the standard random walk MCMC given in formula (1.2) in [1], the two Gaussian-expressions cancel, as do the logarithm and the exponentiation, leaving the simple acceptance described in 1.

This cancellation does not happen when using the pCN-Acceptance probablity. This could explain the poorer performance of pCN when directly sampling a probablity distribution.

1.2.5 Acceptance Probabilities for different MCMC Proposers

Start from Bayes' formula and rewrite the likelyhood $L(u)$ as $\exp(-\Phi(u))$ for a positive scalar function Φ called the potential:

$$\frac{\rho(u)}{\rho_o(u)} \propto \exp(\Phi(u)).$$

Assuming our prior to be a Gaussian ($\mu_0 \sim \mathcal{N}(0, \mathcal{C})$).

Then

$$\rho(u) \propto \exp\left(-\Phi(u) + \frac{1}{2}\left\|C^{-1/2}u\right\|^2\right),$$

since $u^T C^{-1} u = (C^{-1/2}u)^T(C^{-1/2}u) = \langle C^{-1/2}u, C^{-1/2}u\rangle = \left\|C^{-1/2}u\right\|^2$, where in the first equality we used C being symmetric.

This is formula (1.2) in [1] and is used in the acceptance probability for the standard random walk (see also Acceptance Probability for Metropolis-Hastings)

$\mathcal{C}^{-1/2}u$ makes problems in infinite dimensions.

Todo: Why exactly is the second term (from the prior) cancelled when doing pCN?

1.2.6 Different formulations of multivariate Gaussians

Is an RV $\xi \sim \mathcal{N}(0, C)$ distributed the same as $C^{1/2}\xi_0$, with $\xi_0 \sim \mathcal{N}(0, \mathcal{I})$?

From wikipedia: Affine transformation $Y = c + BX$ for $X \sim \mathcal{N}(\mu, \Sigma)$ is also a Gaussian $Y \sim \mathcal{N}\left(c + B\mu, B\Sigma B^T\right)$. In our case $X \sim \mathcal{N}(0, \mathcal{I})$, so

$Y \sim \mathcal{N}\left(0, C^{1/2} \mathcal{I} C^{1/2^T}\right) = \mathcal{N}(0, C)$, since the covariance matrix is positive definite, which means it's square root is also positive definite and thus symmetric.

On second thought, it also follows straight from the definition:

$$\mathbf{X} \sim \mathcal{N}(\mu, \Sigma) \Leftrightarrow \exists \mu \in \mathbb{R}^k, A \in \mathbb{R}^{k \times l} \text{ s.t. } \mathbf{X} = \mu + A\mathbf{Z} \text{ with } \mathbf{Z}_n \sim \mathcal{N}(0, 1) \text{ i.i.d}$$

where $\Sigma = AA^T$.

2   Framework/Package Structure

The framework is designed to support an easy use case:

```
proposer = StandardRWProposer(beta=0.25, dims=1)
accepter = AnalyticAccepter(my_distribution)
rng = np.random.default_rng(42)
sampler = MCMCSampler(rw_proposer, accepter, rng)

samples = sampler.run(x_0=0, n_samples=1000)
```

There is only one source of randomness, shared among all classes and supplied by the user. This facilitates reproducability.

Tests are done with `pytest`.

2.1   Distributions

A class for implementing probability distributions.

```
class DistributionBase(ABC):
    @abstractmethod
    def sample(self, rng):
        """Return a point sampled from this distribution"""
        ...
```

The most important realisation is the `GaussianDistribution`, used in the proposers.

```
class GaussianDistribution(DistributionBase):
    def __init__(self, mean=0, covariance=1):
        ...

    def sample(self, rng):
```

```
        ...

    def apply_covariance(self, x):
        ...

    def apply_sqrt_covariance(self, x):
        ...

    def apply_precision(self, x):
        ...

    def apply_sqrt_precision(self, x):
        ...
```

The design of this class is based on the implementation in muq2. The `precision` / `sqrt_precision` is implemented through a Cholesky decomposition, computed in the constructor. This makes applying them pretty fast $(\mathcal{O}(n^2))$.

At the moment the there is one class for both scalar and multivariate Gaussians. This introduces some overhead as it has to work with both `float` and `np.array`. Maybe two seperate classes would be better.

Also, maybe there is a need to implement a Gaussian using the Karhunen-Loéve-Expansion?

2.2 Potentials

A class for implementing the potential resulting from rewriting the likelihood as
$$\mathrm{L}(u) = \exp(-\Phi(u)).$$

```
class PotentialBase(ABC):
    """
    Potential used to express the likelihood;
    d mu(u; y) / d mu_0(u) \propto L(u; y)
    Write L(u; y) as exp(-potential(u; y))
    """
    @abstractmethod
    def __call__(self, u):
        ...
```

```python
@abstractmethod
def exp_minus_potential(self, u):
    ...
```

The two functions return $\Phi(u)$ and $\exp(-\Phi(u))$ respectively. Depending on the concrete potential, one or the other is easier to compute.

Potentials are used in the accepters to decide the relative weight of different configurations. They use the `__call__`-method to do that. Especially for high-dimensional error-terms, the pdf of the error term can become very small, so it is important to implement this computing the logpdf directly instead of manually exponentiating and running into issues with floating point limitations.

2.2.1 AnalyticPotential

This potential is used when sampling from an analytically computable probability distribution, i.e. a known posterior. In this case

$$\exp(-\Phi(u)) = \frac{\rho(u)}{\rho_0(u)},$$

see `theory.org`

2.2.2 EvolutionPotential

This potential results when sampling from the model-equation

$$y = \mathcal{G}(u) + \eta,$$

with $\eta \sim \rho$. The resulting potential can be computed as

$$\exp(-\Phi(u)) = \rho(y - \mathcal{G}(u)).$$

2.3 Proposers

Propose a new state v based on the current one u.

```python
class ProposerBase(ABC):
    @abstractmethod
    def __call__(self, u, rng):
        ...
```

9

2.3.1 StandardRWProposer

Propose a new state as
$$v = u + \sqrt{2\delta}\xi,$$
with either $\xi \sim \mathcal{N}(0, \mathcal{I})$ or $\xi \sim \mathcal{N}(0, \mathcal{C})$ (see section 4.2 in [1]).

This leads to a well-defined algorithm in finite dimensions. This is not the case when working on functions (as described in section 6.3 in [1])

2.3.2 pCNProposer

Propose a new state as
$$v = \sqrt{1 - \beta^2}u + \beta\xi,$$
with $\xi \sim \mathcal{N}(0, \mathcal{C})$ and $\beta = \frac{8\delta}{(2+\delta)^2} \in [0, 1]$ (see formula (4.8) in [1]).

This approach leads to an improved algorithm (quicker decorrelation in finite dimensions, nicer properties for infinite dimensions)(see sections 6.2 + 6.3 in [1]).

The wikipedia-article on the Cholesky-factorization mentions the use-case of obtaining a correlated sample from an uncorrelated one by the Cholesky-factor. This is not implemented here.

2.4 Accepters

Given a current state u and a proposed state v, decide if the new state is accepted or rejected.

For sampling from a distribution $P(x)$, the acceptance probability for a symmetric proposal is $a = \min\{1, \frac{P(v)}{P(u)}\}$ (see `theory.org`)

```python
class ProbabilisticAccepter(AccepterBase):
    def __call__(self, u, v, rng):
        """Return True if v is accepted"""
        a = self.accept_probability(u, v)
        return a > rng.random()

    @abstractmethod
    def accept_probability(self, u, v):
        ...
```

2.4.1 AnalyticAccepter

Used when there is an analytic expression of the desired distribution.

```python
class AnalyticAccepter(ProbabilisticAccepter):
    def accept_probability(self, u, v):
        return self.rho(v) / self.rho(u)
```

2.4.2 StandardRWAccepter

Based on formula (1.2) in [1]:

$$a = \min\{1, \exp(I(u) - I(v))\},$$

with

$$I(u) = \Phi(u) + \frac{1}{2}\left\|\mathcal{C}^{-1/2}u\right\|^2$$

.

See also `theory.org`.

2.4.3 pCNAccepter

Works together with the pCNProposer to achieve the simpler expression for the acceptance

$$a = \min\{1, \exp(\Phi(u) - \Phi(v))\}.$$

2.4.4 CountedAccepter

Stores and forwards calls to an "actual" accepter. Counts calls and accepts and is used for calculating the acceptance ratio.

2.5 Sampler

The structure of the sampler is quite simple, since it can rely heavily on the functionality provided by the Proposers and Accepters.

```python
class MCMCSampler:
    def __init__(self, proposal, acceptance, rng):
        ...

    def run(self, u_0, n_samples, burn_in=1000, sample_interval=200):
        ...

    def _step(self, u, rng):
        ...
```

3 Results

3.1 Analytic sampling from a bimodal Gaussian

3.1.1 Setup

Attempting to recreate the "Computational Illustration" from [1]. They use, among other algorithms, pCN to sample from a 1-D bimodal Gaussian

$$\rho \propto (\mathcal{N}(3,1) + \mathcal{N}(-3,1))\mathbb{1}_{[-10,10]}.$$

Since the density estimation framework for a known distribution is not quite clear to me from the paper, I don't expect to perfectly replicate their results.

They use a formulation of the prior based on the Karhunen-Loéve Expansion that doesn't make sense to me in the 1-D setting (how do I sum infinite eigenfunctions of a scalar?).

The potential for density estimation described in section is also not clear to me (maybe for a similar reason? What is u in the density estimate case?).

I ended up using a normal $\mathcal{N}(0,1)$ as a prior and the potential described before, and compared the following samplers:

- (1) `StandardRWProposer` $(\delta = 0.25)$ + `AnalyticAccepter`

- (2) `StandardRWProposer` $(\delta = 0.25)$ + `StandardRWAccepter`

- (3) `pCNProposer` $(\beta = 0.25)$ + `pCNAccepter`

The code is in `analytic.py`.

3.1.2 Result

All three samplers are able to reproduce the target density 1

Figure 1: Burn-in: 1000, sample-interval: 200, samples: 500

The autocorrelation decays for all samplers: 2. However, the pCN doens't do nearly as well as expected. This could be the consequence of the awkward formulation of the potential or a bad prior.

12

Figure 2: AC of bimodal distribution. pCN takes forever to decorrelate

3.2 Bayesian inverse problem for $\mathcal{G}(u) = \langle g, u \rangle$

For $\mathcal{G}(u) = \langle g, u \rangle$ the resulting posterior under a Gaussian prior is again a Gaussian. The model equation is

$$y = \mathcal{G}(u) + \eta$$

with:

- $y \in \mathbb{R}$

- $u \in \mathbb{R}^n$

- $\eta \sim \mathcal{N}\left(0, \gamma^2\right)$ for $\gamma \in \mathbb{R}$

A concrete realization with scalar u:

- $u = 2$

- $g = 3$

- $\gamma = 0.5$

13

- $y = 6.172$

- prior $\mathcal{N}\left(0, \Sigma_0 = 1\right)$

leads to a posterior with mean $\mu = \frac{(\Sigma_0 g)y}{\gamma^2 + \langle g, \Sigma_0 g \rangle} \approx 2$, which is what we see when we plot the result 3. The pCN-Sampler with $\beta = 0.25$ had an acceptance rate of 0.567.


Figure 3: $N = 5000, \mu \approx 2$

For $n > 2$, the resulting posterior can not be plotted anymore. However, it is still Gaussian with given mean & covariance. Can just compare the analytical values to the sample values. Verify that the error decays like $\frac{1}{\sqrt{N}}$.

3.3 Bayesian inverse problem for $\mathcal{G}\left(u\right) = g(u + \beta u^3)$

Since the observation operator is not linear anymore, the resulting posterior is not Gaussian in general. However, since the dimension of the input u is 1, it can still be plotted.

The concrete realization with:

- $g = [3, 1]$

- $u = 0.5$

- $\beta = 0$

- $y = [1.672, 0.91]$

- $\gamma = 0.5$

- $\eta \sim \mathcal{N}\left(0, \gamma^2 I\right)$

- prior $\mathcal{N}\left(0, \Sigma_0 = 1\right)$

however leads to a Gaussian thanks to $\beta = 0$. The mean is $\mu = \frac{\langle g, y \rangle}{\gamma^2 + |g|^2} \approx$ 0.58. Plot: 4

The pCN-Sampler with $\beta = 0.25$ (different beta) had an acceptance rate of 0.576.

Figure 4: $\quad N = 5000, \mu \approx 0.58$

For $\beta \neq 0$, the resulting posterior is not a Gaussian. Still $n = 1$, so it can be plotted. Just numerically normalize the analytical expression of the posterior?

3.4 Lorenz96 model

3.4.1 Model

Based on: Properly cite this!

Lorenz, E. N. (1996). Predictability—A problem partly solved. In Reprinted in T. N. Palmer & R. Hagedorn (Eds.), Proceedings Seminar on Predictability, Predictability of Weather and Climate, Cambridge UP (2006) (Vol. 1, pp. 1–18). Reading, Berkshire, UK: ECMWF.

Equation A system of ODEs, representing the coupling between slow variables X and fast, subgrid variables Y. The system is used in [2] to illustrate different algorithms for earth system modelling.

$$\frac{\mathrm{d}X_k}{\mathrm{d}t} = -X_{k-1}(X_{k-2} - X_{k+1}) - X_k + F - hc\bar{Y}_k \tag{2}$$

$$\frac{1}{c}\frac{\mathrm{d}Y_{j,k}}{\mathrm{d}t} = -bY_{j+1,k}(Y_{j+2,k} - Y_{j-1,k}) - Y_{j,k} + \frac{h}{J}X_k \tag{3}$$

- $X = [X_0, ..., X_{K-1}] \in \mathbb{R}^K$

- $Y = [Y_{j,0}|...|Y_{j,K-1}] \in \mathbb{R}^{J \times K}$
  $Y_{j,k} = [Y_{0,k}, ..., Y_{J-1,k}] \in \mathbb{R}^J$

- $\bar{Y}_k = \frac{1}{J}\sum_j Y_{j,k}$

- periodic: $X_K = X_0$, $Y_{J,k} = Y_{0,k}$

- Parameters $\Theta = [F, h, c, b]$

- h: coupling strength

- c: relative damping

- F: external forcing of the slow variables (large scale forcing)

- b: scale of non-linear interaction of fast variables

- $t = 1 \Leftrightarrow 1$ day (simulation duration is given in days)

b **or** J? In the original paper, the equations are given in a different form, namely all explicit occurences of J above (in the fast-slow interaction) are replaced by b. Since in both concrete realizations (1996 & 2017) are identical and conviniently have $b = J = 10$, the difference doesn't lead to different results for that setup.

"Looking ahead" vs. "Looking back" Comparing nonlinearity terms

$$-X_{k-1}(X_{k-2} - X_{k+1})$$

$$-bY_{j+1,k}(Y_{j+2,k} - Y_{j-1,k})$$

for a given Y_k, does the "direction" of the $Z_{k\pm1}Z_{k\pm2}$ $(Z = X, Y)$ matter?

I don't think so, since the interaction with the other variable is only via point-value and average, and the nonlinearity is periodic.

A bit more formally: The PDE is invariant under "reversing" of the numbering: $Y_{j,k} \to Y_{J-j,k}$ which is the same as switching $+ \leftrightarrow -$ in the only "asymmetric" term.

Addendum 2 days later: Need to define more clearly what it means for the direction to <u>matter</u>. In the original paper on page 12, it is described how "active areas [..] propagate slowly eastward", while "convective activity tends to propagate westward within the active areas" (rephrased from paper). The paper also explicitly mentions the signs of the subscripts in that context. So some characteristics of the solution are definitely affected. What about the stuff we care about (statistical properties, chaotic behaviour)?

Properties For $K = 36$, $J = 10$ and $\Theta = [F, h, c, b] = [10, 1, 10, 10]$ there is chaotic behaviour.

Energy The nonlinearities conserve the energies within a subsystem:

- $E_X = \sum_k X_k^2$

  $\frac{1}{2}\frac{\mathrm{d}(\sum_k X_k^2)}{\mathrm{d}t} = \sum_k X_k \frac{\mathrm{d}X_k}{\mathrm{d}t} = -\sum_k (X_k X_{k-1} X_{k-2} - X_{k-1} X_k X_{k+1}) = 0$,

  where the last equality follows from telescoping + periodicity

- $E_{Y_k} = \sum_j Y_{j,k}^2$

  which follows analogously to the X -case

The interaction between fast and slow variables conserves the total energy:

- $E_T = \sum_k (X_k^2 + \sum_j Y_{j,k}^2)$

  $\frac{1}{2}\frac{\mathrm{d}E_T}{\mathrm{d}t} = \sum_k X_k \frac{\mathrm{d}X_k}{\mathrm{d}t} + \sum_j Y_{j,k} \frac{\mathrm{d}Y_{j,k}}{\mathrm{d}t} = \sum_k X_k(-\frac{hc}{J}\sum_j Y_{j,k}) + \sum_j Y_{j,k}(\frac{hc}{J}X_k) = \sum_k -\frac{hc}{J}X_k(\sum_j Y_{j,k} + \frac{hc}{J}X_k(\sum_j Y_{j,k})) = 0$

In the statistical steady state, the external forcing F (as long as its positive) balances the dampling of the linear terms.

Averaged quantities

$$\langle \phi \rangle = \frac{1}{T} \int_{t_0}^{t_0+T} \phi(t) \, \mathrm{d}t$$

(or a sum over discrete values)

Long-term time-average in the statistical steady state: $\langle \cdot \rangle_\infty$

-
$$\langle X_k^2 \rangle_\infty = F \langle X_k \rangle_\infty - hc \langle X_k \bar{Y}_k \rangle_\infty \; \forall k \qquad (4)$$

  (multiply X-equation by X, all X_k s are statistically equivalent, $\frac{\mathrm{d}\langle X \rangle}{\mathrm{d}t} = 0$ in steady state)

-
$$\left\langle \bar{Y}_k^2 \right\rangle_\infty = \frac{h}{J} \langle X_k \bar{Y}_k \rangle_\infty \; \forall k \qquad (5)$$

(Quasi) Ergodicity Whether chaotic regions of the phase space of a system are ergodic seems not be an easy question to answer (citation needed probably) [1] Are there any inaccessible regions in phase space for the Lorenz system? I can't think of any. However, there seem to be "traps" that take the system out of it's chaotic behaviour ($X_i = c$, $Y_i = a$). These somehow destroy ergodicity. Are they somehow "measure 0" or something?)/. However, for the purposes of this section (which deals with finite time anyway), it is enough to assert that

for the Lorenz system, for sufficiently long times, the time-average converges to the "space-average" over phase-space:

$$\langle f \rangle_\infty = \lim_{T \to \infty} \int_0^T f(Z(t)) \, \mathrm{d}t = \int_{\mathbb{R}^{K(J+1)}} f(x)\rho(x) \, \mathrm{d}x \qquad (6)$$

where $Z(t)$ is a phase space trajectory of the system and $\rho(x)$ is the probaility of the system in the statistical steady state to be in state x.

One sufficiently long simulation of the system gives information about all accessible [2] initial conditions. As a consequence, as long as the integration time of the system is "long enough", the chosen initial condition is meaningless and can even vary without changing the behaviour of the observation operator.

---

[1] Are there any inaccessible regions in phase space for the Lorenz system? I can't think of any. However, there seem to be "traps" that take the system out of it's chaotic behaviour ($X_i = c$, $Y_i = a$). These destroy ergodicity. Are they somehow "measure 0" or something?

[2] Here a more precise definition of ergodicity of the system would help out. What I mean is "all sensible initial conditions".

3.4.2 Model implementation

Implementing the model in python and using a locally 5-th order RK solver yields the following results (inital conditions are just uniformly random numbers in $[0, 1)$ since they don't matter for the long-term evolution of the chaotic system):

Reproducting the results of the original paper Running the setup with $K = 36, J = 10, (F, h, c, b) = (10, 1, 10, 10)$ gives the following states 5, which qualitatively agree with the results from Lorenz.


Figure 5: System around $T = 20$

The decay of the linear term and the forcing of the slow variables balance out after reaching the steady state, however there is a much bigger fluctuation in the energy than expected 6.

Equilibrium averages Analysis suggests certain long-term averages to be equal in the equilibrium.

3.4.3 MCMC

General point: The `RK45` method uses a predictor/corrector step and thus does non-uniform timesteps. However, in the following I compute time-averages with a simple `np.mean`, ignoring the different length of timesteps. It would be not impossible to write my own `time_average(y, t)`-function that takes the non-uniform timesteps into account. However I'm not sure how

Figure 6: Energies in the system. $E_X >> E_{Y_k} > 0$

Figure 7: RMSE for long-term averages 4 and 5. Averaged over 10 runs

necessary this is, considering a forward-integration takes (2000) timesteps, so I suspect that differences are washed out a bit?

Setup Denote the Lorenz-96 system 2, 3 with parameters $\tilde{u} = [F, h, c, b]$ as $\mathcal{M}[\tilde{u}]$. It acts on the initial condition $z_0 = [X_0, Y_0] \in \mathbb{R}^{K(J+1)}$ to evolve the system for N_t timesteps and generate the phase space trajectory $Z = [\begin{smallmatrix} X_1 \\ Y_1 \end{smallmatrix} | \cdots | \begin{smallmatrix} X_{N_t} \\ Y_{N_t} \end{smallmatrix}] \in \mathbb{R}^{K(J+1) \times N_t}$:

$$Z = \mathcal{M}[\tilde{u}] z_0$$

Define the "moment function" $f(z) : \mathbb{R}^{K(J+1)} \to \mathbb{R}^{5K}$

$$f(z) = \begin{bmatrix} \bar{X} \\ \bar{Y} \\ \overline{X^2} \\ \overline{X\bar{Y}} \\ \overline{Y^2} \end{bmatrix} \tag{7}$$

The MCMC-Algorithm then samples based on:

$$\langle f \rangle_\infty = \langle f \rangle_T (u) + \eta$$

with:

- $\langle f \rangle_\infty \approx \langle f \rangle_{T_r}$ with $T_r >> T$ over a simulation $\mathcal{M}[u^*] z_0$

- $\langle f \rangle_T (u)$ the time average over a simulation $\mathcal{M}[u_p + u] z_0$

- Due to the ergodic properties of $\mathcal{M}$ 3.4.1 , it doesn't really matter what z_0 is

- The parameter vector comes in many different variations:

    - $u^* \in \mathbb{R}^4_{\geq 0}$: true underlying parameters, used to compute the "data"

    - $u_p \in \mathbb{R}^4_{\geq 0}$: mean of the prior

    - $u \in \mathbb{R}^4$: pertubations to the prior mean, the actual input to the observation operator

- $\eta \sim \mathcal{N}(0, \Sigma)$, where $\Sigma = r^2 [\text{var}(f)_{T_r}]$, where $r \in \mathbb{R}$ is the "noise level"

1. The parameter vector u

   The theoretical background assumes the prior to be a centered Gaussian ($\mu = 0$). Specifically, it matters during the proposal step, where the step is taken either with scaled sample from the prior or from a centered Gaussian with the covariance of the prior. A compromise would be to just ignore a nonzero prior-mean in the proposal, however I'm not sure if such a prior has other effects that invalidate the algorithm.

2. "Noise level"

   r is scaling of covariance matrix of noise term. This in turn is just step-width in proposal.

   TODO: Verify by checking acceptance rate for different noise levels

Concrete parameters The MCMC-Simulation was carried out with the following parameters:

- $K = 6, J = 4$

- Reference Simulation to get $\langle f \rangle_\infty$ and Σ:

  - $u^* = [F^*, h^*, c^*, b^*] = [10, 10, 1, 10]$
  - $T_r = 500$

- Noise $\eta \sim \mathcal{N}(0, \Sigma)$ with $\Sigma = r^2 \mathrm{diag}(\mathrm{var}(f_{T_r})) \in \mathbb{R}^{5K \times 5K}$

- Noise level $r = 0.5$

- Prior $\mathcal{N}(u_p, \Sigma_0)$

  - $u_p = [F_0, h_0, b_0] = [12, 8, 9]$
  - $\Sigma_0 = \mathrm{diag}([10, 1, 10])$
  - c was excluded from the sampling since it is very hard to estimate ("bad statistics")
  - The prior was chosen closer to the true value to make the job of the algorithm easier

- Sampling with pCN proposer and acccepter with $\beta = 0.25$

  - Evaluating the observation operator with a model-simulation of $T = 20$

- Start sampling very close to true value: $u_0 = [-1.9, 1.9, 0.9]$ so that $u^* \approx u_p + u_0$

- This means we can use a short burn-in of 100

- Sample $N = 2500$ with a sample-interval of 2

- The sample interval of 2 is very short, especially considering the long correlation time see below. But 2 is also what they used in the ESM paper.

Result

Density plot for posterior   The resulting density plots show a improvement from the prior towards the true value 8. The estimation of the parameter F seems to be easier than b, where the prior and the posterior seem pretty much identical.

This slight improvement is however not unexpected, as the simulations I've done are much shorter than the ones in [2] $(K, J) = (6, 4)$ vs $(36, 10)$, $T = 20$ vs 100, $T_r = 500$ vs 46,416)

Should I do some more analysis here, like reporting sample means and covariances to compare posterior/prior not just visually?
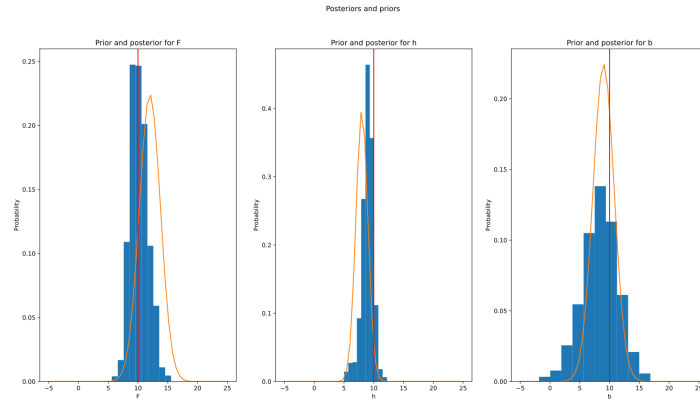

Figure 8:   Prior and Posteriors after a 5000 sample MCMC run

ACF   While the plot of the F-values looks pretty uncorrelated to me 10, the ACF functions decays much slower than expected 9. I there maybe something wrong with how I compute it?
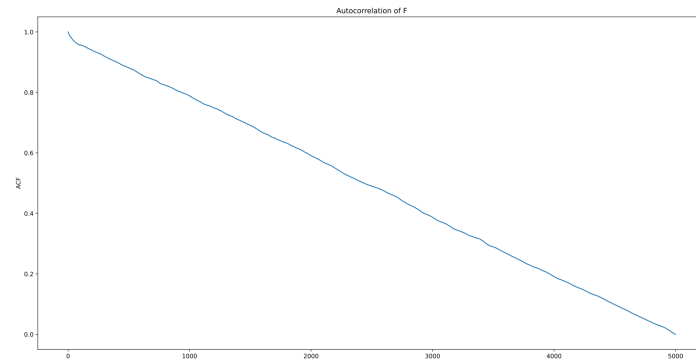
Figure 9: Autocorrelation of F during the MCMC sampling

Figure 10: Value of F during the MCMC sampling

References

[1] S. L. Cotter, G. O. Roberts, A. M. Stuart, and D. White. MCMC Methods for Functions: Modifying Old Algorithms to Make Them Faster. *Statistical Science*, 28(3):424–446, August 2013. Publisher: Institute of Mathematical Statistics.

[2] Tapio Schneider, Shiwei Lan, Andrew Stuart, and João Teixeira. Earth System Modeling 2.0: A Blueprint for Models That Learn From Observations and Targeted High-Resolution Simulations. *Geophysical Research Letters*, 44(24):12,396–12,417, 2017. _eprint: https://agupubs.onlinelibrary.wiley.com/doi/pdf/10.1002/2017GL076101.

[3] A. M. Stuart. Inverse problems: A Bayesian perspective. *Acta Numerica*, 19:451–559, May 2010.